

F25-314-D-CoWriteIA

Project Team

Aisha Siddiqa	22I-1281
Ayesha Ejaz	22I-0899
Junaid Asrar	22I-0770

Session 2022-2026

Supervised by

Dr. Ali Zeeshan Ijaz



Department of Computer Science

**National University of Computer and Emerging Sciences
Islamabad, Pakistan**

June, 2026

Contents

1	Introduction	1
1.1	Existing Solutions	1
1.2	Problem Statement	2
1.3	Scope	2
1.4	Modules	3
1.4.1	Module 1: Project Indexing and Semantic Memory	3
1.4.2	Module 2: Context-Aware Writing Assistant	3
1.4.3	Module 3: Character and Scene Management	4
1.4.4	Module 4: Dialogue Generation	4
1.4.5	Module 5: Research Integration	4
1.4.6	Module 6: Project Query Interface	4
1.4.7	Module 7: Gap Analysis Module	4
1.5	Work Division	5
2	Project Requirements	7
2.1	Use-case	7
2.2	Event–Response Table	8
2.3	Functional Requirements	9
2.3.1	Module 1: Project Indexing and Semantic Memory	9
2.3.2	Module 2: Context-Aware Writing Assistant	9
2.3.3	Module 3: Character and Scene Management	9
2.3.4	Module 4: Dialogue Generation	9
2.3.5	Module 5: Research Integration Module	10
2.3.6	Module 6: Project Query Interface	10
2.3.7	Module 7: Gap Analysis Module	10
2.4	Non-Functional Requirements	10
2.4.1	Usability	10
2.4.2	Performance	11
2.4.3	Security	11
2.4.4	Reliability	11
2.4.5	Maintainability	11
2.4.6	Compatibility	11

2.4.7	Scalability	11
3	System Overview	13
3.1	Architectural Design	13
3.2	Design Models	15
3.2.1	Sequence Diagrams	15
3.2.2	State Transition Diagram	21
3.3	Domain Model	21
3.3.1	Relationships and Associations	23
3.3.2	Domain Model Diagram	23
3.4	Data Design	25
3.4.1	Class Diagram	26
4	Implementation and Testing	29
4.1	Algorithm Design	29
4.1.1	High-level Flow	30
4.1.2	Pseudocode: Document Chunking and Preprocessing	31
4.1.3	Pseudocode: Embedding Generation and Storage	33
4.1.4	Pseudocode: Context-aware Retrieval (RAG)	33
4.1.5	Pseudocode: AI Response Generation (RAG + LLM Call)	34
4.2	External APIs and SDKs	37
4.3	Testing Details	37
4.3.1	Black Box Testing	37
4.3.1.1	Purpose of Black Box Testing	37
4.3.1.2	Testing Environment	37
4.3.1.3	Functional API Coverage	38
4.3.1.4	Workflow and Scenario Testing	38
4.3.1.5	Negative and Security Testing	38
4.3.1.6	Error and Edge Case Validation	38
4.3.1.7	Black Box Test Evidence	38
4.3.1.8	Black Box Testing Summary	40
4.3.2	Unit Testing	40
4.3.2.1	Implemented Unit Test Coverage	40
4.3.2.2	Testing Approach	40
4.3.2.3	Failure and Boundary Validation	40
4.3.2.4	Authentication and File Services Unit Tests	41
4.3.2.5	Search, Text Extraction, and Embedding Unit Tests	41
4.3.2.6	LLM, Editing, and Copilot Unit Tests	41
4.3.2.7	RAG Context, Entity Extraction, and Relationship Discovery Unit Tests	42
4.3.2.8	Unit Test Summary	42

4.3.2.9	Current Iteration Test Additions	42
4.3.3	Test Evidence	44
4.4	Test Summary	45
4.4.1	Overall Testing Statistics	45
4.4.2	Integration Testing	48
4.4.2.1	Summary	48
4.4.2.2	Frontend Testing Infrastructure	48
4.5	Current Iteration Features	49
4.5.1	Dialogue Generation Module	49
4.5.1.1	Module Summary	49
4.5.1.2	Algorithm: Context-Aware Dialogue Generation	50
4.5.2	Character Repository Module	52
4.5.2.1	Module Summary	52
4.5.2.2	Algorithm: Character Profile Extraction and Update	53
4.5.3	Research Integration Agent Module	55
4.5.3.1	Module Summary	55
4.5.3.2	Algorithm: Web Research and AI Synthesis	55
5	Deployment and Production Infrastructure	59
5.1	Deployment Architecture Overview	59
5.2	Frontend Deployment on Vercel	60
5.2.1	Build Configuration	60
5.2.2	Vercel Project Setup	60
5.2.3	API Routing and CORS	61
5.2.4	Preview Deployments	61
5.3	Backend API Deployment	61
5.3.1	Runtime and Server	61
5.3.2	Startup Lifecycle	62
5.3.3	Middleware Stack	62
5.3.4	Environment Configuration	62
5.4	MongoDB Atlas Integration	63
5.4.1	Database Design and Collections	63
5.4.2	Atlas Configuration	64
5.4.3	Vector Search Index	64
5.4.4	Connection Pooling	65
5.5	Vector Database: ChromaDB	65
5.5.1	Role in the System	65
5.5.2	Deployment	65
5.5.3	Collection Management	66
5.5.4	Fallback Strategy	66

5.6	Redis: Caching and Task Queue	66
5.6.1	Session and Response Caching	66
5.6.2	Celery Task Queue	66
5.6.3	Redis Deployment Options	67
5.7	API Communication and Integration	67
5.7.1	Frontend-to-Backend API Calls	67
5.7.2	Third-Party API Integrations	68
5.7.3	Request Flow: End-to-End Example	68
5.8	Monitoring and Observability	69
5.8.1	Application-Level Logging	69
5.8.2	Health Check Endpoint	69
5.8.3	Error Tracking	69
5.8.4	Performance Metrics	69
5.8.5	Vercel Analytics	70
5.9	Security Considerations	70
5.9.1	Authentication and Authorisation	70
5.9.2	Secret Management	70
5.9.3	Input Validation	71
5.9.4	HTTPS Enforcement	71
5.10	Deployment Summary	71
6	Conclusions and Future Work	73
6.1	Conclusion	73
6.1.1	What Was Built	73
6.1.2	Testing and Validation	74
6.1.3	Deployment	74
6.1.4	Reflection on Objectives	74
6.2	Future Work	75
6.2.1	Retrieval Quality and Scalability	75
6.2.2	Evaluation Framework	75
6.2.3	Model and Provider Flexibility	76
6.2.4	Collaborative Writing	76
6.2.5	Interface and Accessibility	76
6.2.6	Infrastructure Hardening	76
	References	77
A	Appendices	78
A.1	Appendix A: System Diagrams	78
A.1.1	Use Case Diagram	78
A.1.2	Activity Diagram	79

A.2	Appendix B: Architecture Diagram	80
A.2.1	CoWriteIA Detailed Architecture Diagram	80
A.3	Appendix C: User Interface Screenshots	82
A.3.1	Landing Page	82
A.3.2	Dashboard Interface	83
A.3.3	Writing Environment Interface	83
A.3.4	Writing Assistant Interface	84
A.3.5	Login Interface	84
A.4	Appendix D: Detailed Use Case Specification	86
A.4.1	UC-01: Generate Context-Aware Writing	86
A.5	Appendix E: Pseudo-API Snippets (Call Signatures)	87

List of Figures

1.1	Gantt Chart - Project Timeline	6
3.1	Box and Line Diagram	15
3.2	Sequence Diagram: File Upload & Chunking	16
3.3	Sequence Diagram: Embedding Generation & Storage	18
3.4	Sequence Diagram: Semantic Search / RAG Context Retrieval	19
3.5	Sequence Diagram: AI Response Generation	20
3.6	Sequence Diagram: Reprocessing & Incremental Update	21
3.7	State Transition Diagram for CoWriteIA	22
3.8	Domain Model Diagram for CoWriteIA	24
3.9	Entity–Relationship Diagram for CoWriteIA	26
3.10	Class Diagram for CoWriteIA	27
4.1	Black Box Test Evidence 1	38
4.2	Black Box Test Evidence 2	39
4.3	Black Box Test Evidence 3	39
4.4	Black Box Test Evidence 4	39
4.5	Black Box Test Evidence 5	40
4.6	Unit Test Evidence 1	41
4.7	Unit Test Evidence 2	41
4.8	Unit Test Evidence 3	41
4.9	Unit Test Evidence 4	42
4.10	Test Execution Results	45
4.11	Integration API Test Evidence	48
A.1	Use Case Diagram CoWriteIA	78
A.2	Activity Diagram	79
A.3	CoWriteIA Multi-Tier Architecture Diagram	80
A.4	CoWriteIA Landing Page	82
A.5	CoWriteIA Dashboard Interface	83
A.6	Writing Environment Interface	83
A.7	Writing Assistant Interface	84
A.8	Login Interface	84

List of Tables

1.1	Comparison of Existing Solutions	2
1.2	Work Division for Iteration I	5
1.3	Work Division for Iteration II	5
1.4	Work Division for Iteration III	5
1.5	Work Division for Iteration IV	6
2.1	Event–Response Table for CoWriteIA	8
3.1	Architectural Tiers and Responsibilities	14
4.1	External APIs and SDKs Used	37
4.2	Overall Backend Testing Summary	42
4.3	Current Iteration Test Coverage by Feature	44
4.4	Comprehensive Testing Statistics	46
5.1	Production Infrastructure Components	60
5.2	Frontend Environment Variables (Vercel)	61
5.3	Backend Production Environment Variables	63
5.4	MongoDB Collections in Production	64
5.5	Third-Party API Integrations	68
5.6	CoWriteIA Production Deployment Summary	71
A.1	Detailed Use Case: Generate Context-Aware Writing	86

Chapter 1

Introduction

CoWriteIA is an AI-powered writing assistant designed to support creative writers, novelists, researchers, and content creators. Modern writing workflows require managing notes, drafts, characters, scenes, and references across several disconnected tools, which often breaks focus and reduces productivity. Writers struggle to maintain narrative consistency, track earlier ideas, and keep a coherent writing style when working on long projects. Existing tools offer only isolated features such as grammar correction or basic generation and do not provide project-level understanding or semantic memory [5, 8].

CoWriteIA addresses these issues by creating a unified, intelligent workspace where all project files are indexed, searchable, and semantically connected. By integrating project-level memory, context-aware writing, dialogue generation, character management, and research support, the system helps writers maintain consistency and improve their creative process. The platform is designed to reduce cognitive load, avoid fragmented workflows, and provide meaningful AI assistance throughout the writing journey. This chapter presents the background, existing solutions, problem statement, scope, project modules, and work division that ground the implemented system.

1.1 Existing Solutions

Several writing and AI-assisted tools exist, but each focuses on limited aspects of the writing process. Grammarly offers grammar correction and style suggestions but lacks deep contextual awareness across long projects. Notion AI and Jasper AI provide generative assistance but do not maintain story-level continuity or user-specific writing style. Tools like Scrivener help with organization but have no semantic understanding or AI memory. As a result, writers must repeatedly switch between applications to manage notes, drafts, characters, and research, leading to inefficiency and inconsistent writing flow.

Table 1.1: Comparison of Existing Solutions

System Name	System Overview	System Limitations
Grammarly	Provides grammar correction, clarity improvements, and tone suggestions.	No project awareness, no memory of earlier chapters, no semantic search.
Notion AI	Offers AI-assisted content generation and note organization.	Cannot maintain narrative consistency; lacks dialogue and character support.
Scrivener	Strong organizational tool for large writing projects with chapter/scene structure.	No AI support, no semantic retrieval, no automated gap or style analysis.

1.2 Problem Statement

Writers working on long-form projects often lose track of earlier ideas, character traits, plot points, and stylistic decisions. This leads to inconsistencies, repeated ideas, and a break in narrative flow. Existing tools either provide isolated writing support or document organization but do not combine both with meaningful context. AI tools can generate text but fail to maintain project-level continuity, making the generated text feel disconnected from the writer’s established style or storyline. Writers spend significant time searching through older drafts to recall information [3, 5].

CoWriteIA addresses these problems through an intelligent system that continuously indexes project content, retrieves relevant context, and assists writers in generating text aligned with their established narrative and style. By maintaining a unified knowledge base and supporting character consistency, scene management, dialogue generation, and semantic search, the system reduces cognitive overhead and improves creative flow.

1.3 Scope

The scope of **CoWriteIA** covers an intelligent, AI-driven writing environment designed to support long-form creative and research-oriented projects. The system focuses on project-level understanding, semantic reasoning, and context-aware assistance to help writers maintain consistency, creativity, and workflow continuity.

In-Scope Features

- **Project Management and Indexing** — Upload, organize, and store documents across multiple writing projects.
- **Semantic Search and Retrieval** — Retrieve information using a vector-based semantic memory instead of keyword matching [4, 8].
- **Context-Aware Writing Assistance** — Generate or refine text that aligns with

prior chapters, scenes, tone, and writing style.

- **Dialogue and Narrative Support** — Produce character-consistent dialogues and maintain scene-level continuity.
- **Character and Knowledge Storage** — Maintain detailed character profiles, relationships, and narrative facts.
- **Research Integration** — Fetch, summarize, and integrate external factual information to support realistic writing.
- **Style Adaptation** — Adapt AI-generated content to match the user’s writing voice across the project.

Out-of-Scope Features

- Plagiarism detection or originality scoring.
- Multimedia editing such as images, audio, or video.
- Complete publishing or formatting pipelines (e.g., EPUB, print layout).
- Full-fledged grammar-only tools or scriptwriting-specific toolkits.

Scope Summary

CoWriteIA is limited to enhancing the *writing process itself*: maintaining narrative consistency, providing intelligent context-aware support, and simplifying research-driven creative decision making. The system does not aim to replace editing, publishing, or multimedia tools, but instead acts as a unified, intelligent workspace to elevate author productivity and coherence throughout the creative workflow.

1.4 Modules

The project is organized into several modules, each responsible for a specific part of the writing workflow.

1.4.1 Module 1: Project Indexing and Semantic Memory

This module extracts, embeds, and organizes all project content into a semantic database for context-aware retrieval [8].

1. Automatic project indexing and embedding generation.
2. Semantic search based on meaning instead of keywords.

1.4.2 Module 2: Context-Aware Writing Assistant

This module provides intelligent writing suggestions aligned with the user’s tone and project context [5].

1. Generates coherent drafts aligned with past content.
2. Retrieves relevant information to maintain consistency.

1.4.3 Module 3: Character and Scene Management

This module stores and manages characters, scenes, and narrative details.

1. Character profiles with traits and relationships.
2. Scene storage and tracking.

1.4.4 Module 4: Dialogue Generation

This module generates natural, character-consistent dialogue suggestions.

1. Dialogue generation based on personality and context.
2. Supports narrative flow within scenes.

1.4.5 Module 5: Research Integration

This module retrieves factual data from external sources for realistic writing.

1. Web-based research retrieval.
2. Insertable factual references.

1.4.6 Module 6: Project Query Interface

This module allows users to ask natural-language questions about their own project.

1. Semantic question-answering.
2. Source-linked responses.

1.4.7 Module 7: Gap Analysis Module

This module analyzes draft content to find missing or weak areas.

1. Identifies incomplete sections.
2. Suggests improvements for clarity and consistency.

Collectively, these seven modules are implemented and integrated in CoWriteIA, providing a unified workflow for planning, writing, retrieval, assistance, and quality improvement.

1.5 Work Division

Development work was carried out across four major iterations. A summary of responsibilities is presented in the tables below.

Iteration I Tasks

Table 1.2: Work Division for Iteration I

Task	Ayesha	Junaid	Aisha
SRS Document	✓	✓	✓
UML Diagrams	✓	✓	✓
UI Design		✓	✓
Database Connection		✓	✓
Frontend (Main Pages)	✓		
Project Indexing Agent	✓	✓	
Knowledge Retrieval Agent		✓	✓
ChatBot Support		✓	

Iteration II Tasks

Table 1.3: Work Division for Iteration II

Task	Ayesha	Junaid	Aisha
Inline Copilot Support	✓		
Context-Aware Writing Module	✓		
Gap Analysis Module			✓
Style Adaptation		✓	
Testing	✓	✓	✓
Documentation (All Sections)	✓	✓	✓

Iteration III Tasks

Table 1.4: Work Division for Iteration III

Task	Ayesha	Junaid	Aisha
Character Repository	✓		
Research Integration Agent			✓
Dialogue Generation		✓	
Detailed Testing	✓	✓	✓
UI Update	✓	✓	✓

Iteration IV Tasks

Table 1.5: Work Division for Iteration IV			
Task	Ayesha	Junaid	Aisha
Final System Deployment	✓	✓	✓
Regression Testing	✓	✓	✓
Performance Optimization	✓	✓	✓
Final Thesis Documentation	✓	✓	✓
Project Presentation/Demo Prep	✓	✓	✓

Project Timeline




Figure 1.1: Gantt Chart - Project Timeline

Chapter 2

Project Requirements

This chapter defines the requirements that guide CoWriteIA. These requirements were derived from the Proposal, FYP1-report, and the SRS document, and were retained as the baseline for full implementation in FYP-2. They describe expected system behavior, user interaction patterns, and the quality constraints the system must maintain. The requirements are divided into functional and non-functional categories, with additional details about expected user interactions.

2.1 Use-case

CoWriteIA provides writers with AI-assisted features such as project indexing, semantic retrieval, character management, context-aware writing, and dialogue generation. To clearly define how users interact with these features, system behavior is modeled through detailed use cases.

The use cases and requirements in this chapter map to the complete implemented scope across all seven modules defined in Chapter 1.

A use case describes the interaction between the primary actor (the writer) and the system to accomplish a specific task. Each use case outlines the trigger event, preconditions, main workflow steps, and system responses. These models help ensure that system functionality aligns with user expectations and guide the design of interactive features.

In addition to textual use cases, a visual representation in the form of a Use-Case Diagram provides an overview of the major user interactions with CoWriteIA. This includes actions such as uploading documents, retrieving semantic context, generating content with the writing assistant, managing characters, and initiating dialogue generation. Such diagrams help maintain clarity, consistency, and completeness throughout the system design.

A **detailed use case** example for the core functionality **Generate Context-Aware Writing (UC-01)** is provided in Appendix D. This detailed use case outlines the complete interaction flow including preconditions, main flow steps, alternate flows, exceptions, and

postconditions. The use case demonstrates how the system retrieves semantic context, processes user requests through the AI model, and delivers generated content while handling various edge cases such as missing embeddings or service unavailability.

2.2 Event–Response Table

The following table summarizes how CoWriteIA responds to key user-initiated events during writing, project management, and context-aware content generation.

Event	System Response
User uploads a document	System extracts text, generates embeddings, indexes the content in the vector database, and updates the project repository.
User requests semantic search	System retrieves top relevant embeddings, processes them through semantic ranking, and returns contextually relevant text segments.
User requests context-aware writing assistance	System gathers project context, sends the prompt + context to the AI model, generates aligned text, and displays it in the writing panel.
User selects a character for dialogue generation	System fetches character traits and stored personality data, and generates dialogue consistent with the character’s profile.
User updates character information	System stores the updated traits, recalculates related embeddings, and refreshes character-linked semantic memory.
User performs research query	System connects to external sources (via integrated APIs), retrieves factual data, summarizes it, and provides sources if available.
User initiates gap analysis	System scans the document structure, identifies missing or weak sections, highlights inconsistencies, and suggests improvements.
User edits generated content	System updates the project document, regenerates relevant embeddings, and reindexes modified sections.
User views project overview	System displays indexed chapters, character profiles, notes, semantic connections, and recent AI-generated outputs.
AI model becomes unavailable	System displays an error message, logs the failure, and advises the user to retry once service availability is restored.

Table 2.1: Event–Response Table for CoWriteIA

2.3 Functional Requirements

The functional requirements describe the operations the CoWriteIA system must support. These requirements come directly from the features identified in the Proposal and SRS, and serve as the traceable benchmark for implementation and validation.

2.3.1 Module 1: Project Indexing and Semantic Memory

This module handles ingestion, indexing, and semantic storage of project documents.

1. The system shall allow users to upload project files including text documents, chapters, notes, and research material.
2. The system shall extract, segment, and convert uploaded content into embeddings for semantic retrieval.
3. The system shall store embeddings in a vector database.
4. The system shall provide semantic search capabilities based on meaning rather than keyword matching.
5. The system shall return ranked search results with source references.

2.3.2 Module 2: Context-Aware Writing Assistant

This module generates content aligned with user writing style and project context.

1. The system shall analyze previous project content to understand tone, terminology, and style.
2. The system shall allow users to generate context-aware text suggestions based on previous chapters or notes.
3. The system shall retrieve relevant context automatically when generating new content.
4. The system shall maintain style consistency between newly generated and existing content.

2.3.3 Module 3: Character and Scene Management

This module manages characters, their traits, and related narrative structures.

1. The system shall allow users to create, edit, and store character profiles.
2. The system shall store attributes such as personality, relationships, behaviors, and backstory.
3. The system shall allow users to manage scenes and attach characters to scenes.
4. The system shall assist in retrieving character information when generating story text or dialogue.

2.3.4 Module 4: Dialogue Generation

This module generates consistent, character-matching dialogue.

1. The system shall generate dialogue aligned with character personality and tone.
2. The system shall allow users to request dialogue for specific characters or scenes.
3. The system shall ensure continuity between generated dialogue and existing narrative.

2.3.5 Module 5: Research Integration Module

This module retrieves factual information to support realistic writing.

1. The system shall allow users to search for factual references.
2. The system shall fetch external information from trusted research sources.
3. The system shall present direct links to sources when providing research information.

2.3.6 Module 6: Project Query Interface

This module allows users to ask natural-language questions about their project.

1. The system shall process user queries related to characters, scenes, chapters, or events.
2. The system shall fetch relevant information from the semantic memory.
3. The system shall provide answers with linked source passages.

2.3.7 Module 7: Gap Analysis Module

This module identifies missing or weak areas of the writer's draft.

1. The system shall analyze uploaded chapters for missing elements such as incomplete scenes or inconsistent character behavior.
2. The system shall highlight areas needing expansion or clarification.
3. The system shall provide suggestions to improve narrative flow and completeness.

2.4 Non-Functional Requirements

Non-functional requirements ensure that CoWriteIA performs reliably, efficiently, and securely.

2.4.1 Usability

1. The system shall provide a simple and intuitive interface accessible to writers with minimal technical expertise.
2. The system shall allow users to perform core actions such as uploading files, searching, and generating text within no more than three interactions.
3. The UI shall clearly present semantic search results with source references.

2.4.2 Performance

1. The system shall index uploaded documents within 5 seconds for an average chapter-length file.
2. Semantic search results shall appear within 2 seconds of a query.
3. Generated text responses shall be produced within 3–5 seconds depending on context length.

2.4.3 Security

1. User project data shall be stored securely using encrypted connections.
2. Only authenticated users shall access their own documents.
3. The system shall not use project data for training without user permission.

2.4.4 Reliability

1. The system shall maintain uptime of at least 99
2. The system shall recover gracefully from API or model failures by retrying or presenting fallback responses.

2.4.5 Maintainability

1. The codebase shall follow modular architecture to allow updates to individual agents.
2. The system shall support integration of new language models without requiring major structural changes.

2.4.6 Compatibility

1. The system shall run on modern browsers including Chrome, Edge, and Firefox.
2. The frontend shall be built using Next.js and shall be compatible with desktop and tablet interfaces.

2.4.7 Scalability

1. The vector database shall support scaling as the user creates larger projects.
2. The system shall handle multiple concurrent requests without significant performance degradation.

Chapter 3

System Overview

Introduction

This chapter provides a high-level description of the system's overall functionality, context, and architectural design. It explains how the major parts of the implemented system interact and why the system is decomposed into specific modules and tiers.

System Overview

CoWriteIA is an AI-assisted writing platform designed to support writers through various stages of the creative process [2, 6, 7]. The system enables users to:

- Create and manage writing projects
- Upload and organize documents
- Manage character profiles and story elements
- Generate and refine written content using AI assistance

To support these features efficiently and reliably, the system separates concerns into distinct, coordinated components: user interaction, application logic, AI processing, and data management.

3.1 Architectural Design

The architecture of CoWriteIA follows a **multi-tier model**, combining **client-server** and **layered** principles. This structure organizes the system into four logical tiers, each with specific responsibilities to ensure maintainability, scalability, and clarity.

Architectural Tiers

Data Storage Components

- **Main Database:** Stores structured system data (users, projects, characters)
- **Vector Database:** Manages embeddings for semantic retrieval and search

Tier	Component	Responsibilities
Presentation	Frontend Application	User interaction, interface rendering, and user input handling
Application Logic	Backend (API Server)	Authentication, business workflows, and project/document/character management
AI Processing	Processing Worker	Embeddings generation, semantic retrieval, and AI text generation
Data	Storage Systems	Structured data, vector storage, and file management

Table 3.1: Architectural Tiers and Responsibilities

- **File Storage:** Handles uploaded documents and exported files

Architectural Diagram

A detailed architecture diagram illustrating the complete multi-tier structure of CoWriteIA is provided in Appendix B (Figure A.3).

Diagram Development Process

- **Initial Design Stage:** Create a *Box and Line Diagram* for simpler representation of systems
- **Finalization Stage:** After selecting the architecture style/pattern (MVC, Client-Server, Layered, Multi-tiered), create detailed mapping of modules/components to each part of the architecture

Design Principles

This architectural decomposition supports the system’s functional requirements while ensuring:

- **Flexibility:** Components can be modified independently
- **Maintainability:** Clear separation of concerns simplifies updates and debugging
- **Scalability:** Each tier can be scaled independently based on demand
- **Security:** Controlled data flow between tiers with proper authentication

Key Architectural Decisions

- **Separation of AI Processing:** Intensive AI tasks are handled by dedicated workers to maintain API responsiveness
- **Multi-tier Structure:** Enables independent development, testing, and deployment of each tier




Figure 3.1: Box and Line Diagram

- **Modular Design:** Supports controlled extension and feature additions

3.2 Design Models

3.2.1 Sequence Diagrams

This subsection presents sequence diagrams for key system flows, based on the implementation logic described in Chapter 4. Together, they outline interactions between core components across the implemented module set.

1. File Upload & Chunking

Actors: User, API Endpoint, TextExtractionService, TextChunkRepository

1. User uploads file via API endpoint.
2. API calls `TextExtractionService.extract_text_from_file`.
3. API calls `TextExtractionService.chunk_text` to split text into chunks.
4. Chunks are created and saved via `TextChunkRepository`.




Figure 3.2: Sequence Diagram: File Upload & Chunking

2. Embedding Generation & Storage

Actors: API/Orchestrator, EmbeddingService, TextChunkRepository, ChromaService

1. **Batch Encoding:** The system extracts content strings from the chunks and calls `EmbeddingService.generate_embeddings`. This runs a threadpool execution where the model loads lazily to convert text into dense vector representations.
2. **Metadata Persistence:** The system iterates through the original chunks and generated embeddings simultaneously. For each pair, a `TextChunk` model is built and saved to MongoDB via the `TextChunkRepository`.
3. **Vector Indexing:** If the `ChromaService` is configured and available:
 - The system maps the saved chunks to Chroma-compatible entries (ID, vector, and metadata).
 - It performs a bulk insert via `ChromaService.add_embeddings` to enable semantic search.
4. **Fallback:** If Chroma is unavailable, the embeddings remain stored solely within the MongoDB documents.




Figure 3.3: Sequence Diagram: Embedding Generation & Storage

3. Semantic Search / RAG Context Retrieval

Actors: User, API Endpoint, RAGContextService, ChromaService, SearchService, TextChunkRepository

1. User sends a search/query.
2. API calls `RAGContextService.assemble_context`.
3. `RAGContextService` generates query embedding.
4. `ChromaService` or `TextChunkRepository` performs vector search.
5. Top chunks and metadata are assembled and returned.




Figure 3.4: Sequence Diagram: Semantic Search / RAG Context Retrieval

4. AI Response Generation

Actors: User, API Endpoint, RAGContextService, AIChatService, GroqService, MessageRepository

1. User sends a message.
2. API calls `RAGContextService.assemble_context`.
3. Context is formatted for LLM.
4. `AIChatService` builds prompt and calls `GroqService` (LLM).
5. LLM response is cleaned and saved via `MessageRepository`.
6. Response returned to user.




Figure 3.5: Sequence Diagram: AI Response Generation

5. Reprocessing & Incremental Update

Actors: User/Admin, TextChunkRepository, FileRepository, TextExtractionService, EmbeddingService

1. File is marked for reprocessing.
2. Existing chunks are deleted.
3. File is marked ready for processing.
4. Text is extracted and chunked again.
5. Chunks are inserted.
6. Embeddings are regenerated and stored.




Figure 3.6: Sequence Diagram: Reprocessing & Incremental Update

3.2.2 State Transition Diagram

Figure 3.7 illustrates state transitions for key entities based on user actions and system events. The diagram shows how entities move between states:

- **Project:** Transitions from Draft → Active → Archived/Deleted based on user actions
- **Document:** Moves through Creating → Editing → Saved → Publishing → Published states
- **Chat Session:** Cycles between Idle → Processing → Waiting for User → Completed
- **Character:** Changes from Draft → Active → Archived as needed

These transitions support robust lifecycle management and data consistency throughout the implemented system.

3.3 Domain Model

The domain model provides a conceptual view of the system's core entities and their relationships, bridging functional requirements and technical implementation. It captures the key business objects, their attributes, and how they interact within CoWriteIA.




Figure 3.7: State Transition Diagram for CoWriteIA

Entities and Attributes

The major entities of the system include:

- **User:** Represents a system user who owns projects and interacts with the AI. Attributes include `user_id`, `name`, `email`, `date_of_birth`, `password_hash`, `role`, and `created_at`.
- **Project:** A workspace created by the user. Attributes include `project_id`, `title`, `created_at`, and `updated_at`.
- **Document:** Represents a chapter or section of a project. Attributes include `document_id`, `title`, `order_num`, `created_at`, and `updated_at`.
- **Character Repository:** A user-owned collection that groups all character profiles.
- **Character:** Represents a story character. Attributes include `character_id`, `name`, `traits`, `history`, `created_at`, and `updated_at`.
- **ProjectCharacter:** An association class linking characters with specific projects. Stores project-specific values such as age and events.
- **Dialogue:** Represents character dialogues. Attributes include `dialogue_id`, `content`, `context`, and `emotion_tag`.
- **ProjectDialogue:** A linking entity that associates dialogues with a project.
- **Chat Session:** Represents an AI chat interaction initiated by the user. Identified by `chat_id`.
- **Message:** Stores an individual message in a chat session. Attributes include `message_id`, `sent_by`, `timestamp`, and `content`.
- **Search Engine:** Represents the semantic search component.
- **Search Result:** Stores contextual search outputs. Attributes include `result_id`,

query, and content.

- **Knowledge Base:** Stores vector embeddings and source references. Attributes include `embedding_id`, `vector_representation`, and `source_ids`.
- **Exported Project:** Represents generated output files such as PDF or Word exports.

3.3.1 Relationships and Associations

The domain model defines the relationships that structure how information flows across the system:

- A **User** owns multiple **Projects**.
- A **Project** contains multiple **Documents**.
- A **User** owns a **Character Repository**, which stores many **Characters**.
- A **Project** includes multiple **Characters** via **ProjectCharacter**.
- A **Character** speaks one or more **Dialogues**.
- A **Project** associates specific dialogues through **ProjectDialogue**.
- A **User** initiates a **Chat Session**.
- A **Chat Session** stores multiple **Messages**.
- The **LLM Chatbot** interacts with the **Semantic Search Engine** and **Knowledge Base** to generate responses.
- The **Search Engine** produces multiple **Search Results**, which may be injected into messages or documents.

These relationships ensure that writing, character creation, semantic search, and AI interactions remain consistent and interconnected across the system.

3.3.2 Domain Model Diagram

Figure 3.8 presents the complete domain model for CoWriteIA. It includes all major entities, attributes, and associations, including weak entities and conceptual relationships.




Figure 3.8: Domain Model Diagram for CoWriteIA

3.4 Data Design

The data design of CoWriteIA transforms functional requirements into structured data models that support **project creation**, **AI-assisted writing**, **semantic search**, **character management**, and **dialogue modelling**.

Data Organization

The system organizes data into three primary categories:

- **Relational Data:** User accounts, projects, documents, characters, and chat sessions stored in a structured database with defined schemas and foreign key relationships
- **Vector Data:** High-dimensional embeddings generated from text content, enabling semantic similarity search and context retrieval
- **Binary Data:** Uploaded files and exported documents stored separately with meta-data linkage

Data Flow and Transformation

Data flows through multiple transformation stages:

- **Input Processing:** User actions validate and create/update structured database records
- **Embedding Generation:** Text content is transformed into vector representations for AI operations
- **Retrieval Augmentation:** Semantic queries fetch relevant context from vector and relational stores
- **Output Generation:** AI responses and exports combine retrieved data with generated content

This design ensures **data integrity**, **efficient retrieval**, and **scalable processing**. Recent work on efficient local models such as Phi-3 highlights practical on-device capabilities for embedding and retrieval [1].

Figure 3.9 provides the Entity–Relationship Diagram showing database schema structure, while Figure 3.8 presents the conceptual domain model.




Figure 3.9: Entity–Relationship Diagram for CoWriteIA

3.4.1 Class Diagram

The class diagram represents the object-oriented view of the system and shows classes, attributes, and associations used by the application logic. It provides a structural foundation for backend implementation and helps ensure consistency between the conceptual design and the code-level architecture.

Class Diagram




Figure 3.10: Class Diagram for CoWriteIA

Chapter 4

Implementation and Testing

This chapter describes the current implementation status of the CoWriteIA system and the testing strategies applied to validate its functionality. The focus of this phase was to implement the core backend services, major APIs, and conduct rigorous black box testing to ensure that all user-facing functionalities behave correctly according to the system requirements.

The implementation described in this chapter reflects the progress achieved up to the current iteration of the project.

4.1 Algorithm Design

The system implements a Retrieval-Augmented Generation (RAG) pipeline to enable intelligent document-based responses. The algorithm focuses only on the core functional requirements: document processing, semantic indexing, context retrieval, and AI response generation.

The CoWriteIA implementation in this repository realizes a semantic indexing + RAG (retrieval-augmented generation) pipeline. The implemented components (and their primary locations in the codebase) are:

- Document chunking and preprocessing
 - `backend/app/services/text_extraction_service.py :: TextExtractionService.chunk_text / _create_text_chunk`
- Embedding generation and vector storage
 - `backend/app/services/embedding_service.py :: EmbeddingService.generate_single_embedding, generate_embeddings, store_text_chunks_with_embeddings`
 - `backend/app/services/chroma_service.py :: ChromaService.add_embeddings, search_similar`
 - `backend/app/repositories/text_chunk_repository.py :: update_embedding, get_chunks_without_embeddings, vector_search / vector_search_chroma`

ma

- Context retrieval (semantic similarity)
 - `backend/app/services/rag_context_service.py :: RAGContextService.assemble_context`
 - `backend/app/services/search_service.py :: SearchService.semantic_search, find_similar_content`
- AI response generation (RAG orchestration and assistant)
 - `backend/app/services/ai_chat_service.py :: AIChatService.chat`
 - `backend/app/services/copilot_service.py :: CopilotService.generate_suggestion`

The following subsections contain concise algorithm descriptions derived from the actual implementation and precise pseudocode that maps to the repository functions.

4.1.1 High-level Flow

1. User uploads file via endpoints in `backend/app/api/v1/endpoints/files.py`.
2. File processing triggers text extraction and chunking:
 - `TextExtractionService.extract_text_from_file(...)` and `TextExtractionService.chunk_text(...)`
 - Chunks are created as `TextChunk` models and persisted via `TextChunkRepository`.
3. Embeddings are generated for chunks:
 - `EmbeddingService.generate_embeddings(...)` / `generate_single_embedding(...)`
 - Embeddings saved both in chunk metadata (MongoDB) and in ChromaDB via `ChromaService` when available.
4. Semantic search / RAG:
 - `RAGContextService.assemble_context(...)` calls embedding/search services to retrieve top-*k* chunks.
 - `SearchService` handles hybrid/fallback logic (ChromaDB preferred, fallback to MongoDB vector search).
5. LLM generation:
 - `AIChatService.chat(...)` or `CopilotService.generate_suggestion(...)` build a prompt from the retrieved context and call the configured LLM (`GroqService` or other).
6. Results returned to the client; metadata and analytics stored as needed.

Example short sequence (mapping code → runtime)

1. User uploads file → API endpoint in `backend/app/api/v1/endpoints/files.py` triggers `process_uploaded_file`.
2. `process_uploaded_file` calls `TextExtractionService.chunk_text` to create chunks.
3. Indexing background task (or immediate synchronous call) calls `EmbeddingService.generate_embeddings` for chunks.

4. `EmbeddingService` stores embeddings via `TextChunkRepository` and `ChromaService`.
5. Later, user sends search/query → `RAGContextService.assemble_context` retrieves top chunks using `ChromaService` or `SearchService`.
6. `AIChatService.chat` builds prompt with context and calls `GroqService` (LLM); response persisted and returned.

4.1.2 Pseudocode: Document Chunking and Preprocessing

The system splits extracted text into sentence-aware chunks within a target size and optionally applies overlap. Each chunk is materialized using the implementation in `backend/app/services/text_extraction_service.py`.

Algorithm 1 ChunkTextAndCreateChunks

Require: *file_id*, *project_id*, *full_text*, *max_chunk_size*, *overlap* (optional), *start_chunk_index* (default 0)

Ensure: List of TextChunk models *chunks*

```
1: sentences ← TextExtractionService._split_into_sentences(full_text)
2: chunks ← []
3: current_chunk_sentences ← []; current_word_count ← 0; chunk_index ← start_chunk_index
4: for each sentence s in sentences do
5:   sentence_words ← count_words(s)
6:   if current_word_count + sentence_words > max_chunk_size and current_chunk_sentences not empty then
7:     chunk_text ← join(current_chunk_sentences, " ")
8:     start_pos ← compute absolute start in full_text
9:     chunk_model ← TextExtractionService._create_text_chunk(file_id,
        project_id, chunk_text, chunk_index, start_pos, current_word_count)
10:    append chunk_model to chunks
11:    chunk_index ← chunk_index + 1
12:    if overlap configured then
13:      current_chunk_sentences ← take last overlap words/sentences of current_chunk_sentences
14:      current_word_count ← count_words(current_chunk_sentences)
15:    else
16:      current_chunk_sentences ← []; current_word_count ← 0
17:    end if
18:  end if
19:  append s to current_chunk_sentences; current_word_count ← current_word_count + sentence_words
20: end for
21: if current_chunk_sentences not empty then
22:   chunk_text ← join(current_chunk_sentences, " ")
23:   start_pos ← compute absolute start in full_text
24:   chunk_model ← TextExtractionService._create_text_chunk(file_id,
        project_id, chunk_text, chunk_index, start_pos, current_word_count)
25:   append chunk_model to chunks
26: end if
27: return chunks
```

Notes:

- Uses sentence-aware splitting and target size with optional overlap per TextExtractionService.
- Returns TextChunk Pydantic models; persistence occurs later (e.g., in store_text).

t_chunks_with_embeddings).

4.1.3 Pseudocode: Embedding Generation and Storage

This stage encodes chunked text into dense vectors and persists both the metadata and vectors. It maps to `EmbeddingService` and `ChromaService` usage in the codebase.

Algorithm 2 GenerateAndPersistEmbeddings

Require: *project_id, file_id, list_of_chunks* (content strings)

Ensure: List of created/updated `TextChunk` records with embeddings

```

1: chunk_texts ← [ chunk.content for chunk in list_of_chunks ]
2: embeddings ← EmbeddingService.generate_embeddings(chunk_texts)           ▷ Runs
   threadpool synchronous encoding; model loads lazily
3: created_chunks ← []
4: for each (chunk_info, embedding) in zip(list_of_chunks, embeddings) do
5:   chunk_model ← build TextChunk( file_id = file_id, project_id = project_id,
   content = chunk_info.content, start_position = ..., end_position = ...,
   chunk_index = chunk_info.chunk_index, word_count = ...,
   embedding_vector = embedding )
6:   saved_chunk ← TextChunkRepository.create(chunk_model)   ▷ Persist metadata in
   MongoDB
7:   append saved_chunk to created_chunks
8: end for
9: if ChromaService available then
10:  chroma_entries ← map created_chunks to {id, embedding, metadata}
11:  ChromaService.add_embeddings(project_id, chroma_entries)
12: else
13:  Optionally store embeddings only within MongoDB fields
14: end if
15: return created_chunks

```

Implementation highlights:

- **Embedding model:** `sentence-transformers/all-MiniLM-L6-v2` (see embedding service and verification scripts).
- **Dual storage:** Metadata in MongoDB `text_chunk` documents; vectors in ChromaDB via `backend/app/services/chroma_service.py`.
- **Batch background tasks:** Celery jobs generate missing embeddings in `backend/app/tasks/indexing.py` (`batch_generate_embeddings_task`).

4.1.4 Pseudocode: Context-aware Retrieval (RAG)

This step assembles a rich context from semantically similar chunks and optional entity/scene/relationship metadata.

Algorithm 3 AssembleRAGContext

Require: query Q , $project_id$, optional $file_id$, $max_chunks\ k$ **Ensure:** $context_payload$: list of chunks, entities, scenes, relationships

```
1: query_embedding  $\leftarrow$  EmbeddingService.generate_single_embedding(Q)
2: if ChromaService available then
3:   vector_hits  $\leftarrow$  ChromaService.search_similar(project_id,
     query_embedding, n_results = k)
4: else
5:   vector_hits  $\leftarrow$  TextChunkRepository.vector_search(project_id,
     query_embedding, limit = k)
6: end if
7: top_chunks  $\leftarrow$  convert  $vector\_hits$  to chunk payloads (id, content, score, file_id,
   positions)
8: if include_entities then
9:   entities  $\leftarrow$  EntityRepository.get_by_project(project_id)
10:  entity_matches  $\leftarrow$  identify entities mentioned in  $top\_chunks$ 
    (EntityExtractionService or string matching)
11: end if
12: if include_scenes then
13:   scenes  $\leftarrow$  SceneRepository.find_by_project(project_id)
14:   mapped_scenes  $\leftarrow$  map chunks to scenes when positions overlap or proximity matches
15: end if
16: if include_relationships then
17:   relationships  $\leftarrow$  RelationshipRepository.find_by_project(project_id)
18: end if
19: context_payload  $\leftarrow$  {"chunks": top_chunks, "entities": entity_matches,
   "scenes": mapped_scenes, "relationships": relationships (filtered)}
20: return context_payload
```

Notes:

- RAGContextService.assemble_context collects chunk-level results, enriches them with entity and scene metadata, and returns a compact context mapping suitable for prompt building (see backend/app/services/rag_context_service.py).
- SearchService implements hybrid logic and fallback ordering (Chroma $\rightarrow$ MongoDB vectorSearch).

4.1.5 Pseudocode: AI Response Generation (RAG + LLM Call)

This step orchestrates context assembly, prompt construction, LLM invocation, and persistence. It corresponds to AIChatService.chat and CopilotService.generate_suggestion.

Algorithm 4 GenerateAIResponseWithContext**Require:** *user_message*, *project_id*, optional *file_id*, *user_id*, *max_context_chunks***Ensure:** *assistant_message*

```

1: Step A: Assemble RAG context
2: context ← RAGContextService.assemble_context(query = user_message,
        project_id = project_id, file_id = file_id, max_chunks = max_context_chunks)
3: Step B: Format context into LLM prompt
4: formatted_context ← RAGContextService.format_context_for_llm(context)
5: system_prompt ← AIChatService.SYSTEM_PROMPT
6: conversation_history ← ConversationRepository.get_latest_messages(
        conversation_id or new)
7: llm_messages ← build_messages_list(system_prompt, conversation_history,
        user_message, formatted_context)    ▷ system message, recent assistant/user turns,
        current user message, appended context
8: Step C: Call LLM
9: assistant_content ← GroqService.generate(messages = llm_messages,
        max_tokens = ..., model = ...)
10: Step D: Postprocess and store
11: cleaned ← clean_text(assistant_content) ▷ e.g., CopilotService._clean_suggestion
12: MessageRepository.save(role = assistant, content = cleaned, context_used = context)
13: update conversation metadata (message count, token usage, timestamps)
14: return cleaned

```

Implementation notes:

- `AIChatService.chat` orchestrates assembling context, formatting prompts, invoking the LLM, and persisting messages.
- `CopilotService.generate_suggestion` targets inline assistance: it infers writing style from RAG results, calls the LLM, then returns a cleaned suggestion and the context used (see `backend/app/services/copilot_service.py`).

Complexity, Reliability, and Practical Concerns

- **Chunking complexity:** $O(n)$ in the number of sentences per document. Overlap increases effective storage but preserves local context for downstream retrieval.
- **Embedding generation:** Dominated by model inference cost per chunk. Batch generation reduces overhead (indexing background tasks `batch_generate_missing_embeddings`).
- **Vector search:** ChromaDB uses efficient ANN (HNSW) with sub-linear query cost; fallback MongoDB vector aggregation may be slower for large collections.
- **Fault tolerance:**
 - Embedding/model calls may fail due to transient network or provider issues; services use `try/catch` with sensible fallbacks (see embedding service and

- verify scripts).
- RAG assembly respects available services and degrades gracefully (ChromaDB optional; code falls back to MongoDB vector search when unavailable).
- **Reindexing & incremental updates:** `TextExtractionService.reprocess_file` deletes existing chunks and marks the file ready for reprocessing; batch background tasks handle any missing embeddings.

Algorithm 5 `ReprocessFileAndIncrementalUpdate`

Require: `file_id, project_id`**Ensure:** File is reindexed and embeddings regenerated as needed

```

1: TextChunkRepository.delete_by_file(file_id)
2: FileRepository.mark_ready_for_processing(file_id)
3: full_text ← TextExtractionService.extract_text_from_file(file_id)
4: chunks ← TextExtractionService.chunk_text(project_id, file_id,
    full_text, max_chunk_size, overlap)
5: TextChunkRepository.bulk_insert(chunks)
6: enqueue batch_generate_missing_embeddings(project_id, file_id)

```

Concrete File References (Key Implementation Points)

- **Text chunking and models**
 - `backend/app/services/text_extraction_service.py :: chunk_text(), _create_text_chunk(), _split_into_sentences()`
 - `backend/app/repositories/text_chunk_repository.py :: create(), get_by_file(), get_chunks_without_embeddings(), update_embedding()`
- **Embeddings and vector store**
 - `backend/app/services/embedding_service.py :: generate_single_embedding(), generate_embeddings(), store_text_chunks_with_embeddings(), find_similar_chunks()`
 - `backend/app/services/chroma_service.py :: get_or_create_collection(), add_embeddings(), search_similar()`
 - `backend/app/tasks/indexing.py :: index_document_task(), batch_generate_embeddings_task()`
- **RAG and search**
 - `backend/app/services/rag_context_service.py :: assemble_context(), format_context_for_llm()`
 - `backend/app/services/search_service.py :: semantic_search(), find_similar_content(), hybrid logic`
- **LLM orchestration and assistants**
 - `backend/app/services/ai_chat_service.py :: chat()`
 - `backend/app/services/copilot_service.py :: generate_suggestion(), gather_story_context()`

See Appendix [A.5](#) for pseudo-API call signatures.

4.2 External APIs and SDKs

The system integrates multiple third-party APIs and SDKs to support AI capabilities, authentication, and data persistence.

API / SDK	Description	Purpose of Usage	Endpoint/ Function
OpenAI API (v1)	Large language model services	Text generation and embeddings	/v1/chat/completions
FastAPI	Python web framework	REST API backend	@app.get(), @app.post()
PostgreSQL	Relational database system	User and project data storage	psycopg2 driver
VectorDB (FAISS/Pinecone)	Vector similarity engine	Semantic search and retrieval	similarity_search()

Table 4.1: External APIs and SDKs Used

4.3 Testing Details

Testing played a critical role in validating the correctness, reliability, and stability of the implemented system. A combination of black box testing and unit testing strategies was adopted to ensure that the system meets its functional requirements.

4.3.1 Black Box Testing

Black box testing was used to validate the external behavior of the system without reference to internal implementation details. The focus was on verifying that system features produced correct outputs when interacting through public interfaces such as API endpoints.

4.3.1.1 Purpose of Black Box Testing

The objective of this testing was to ensure that system functionality aligned with the documented functional requirements by validating response codes, output structures, and observable system behavior.

4.3.1.2 Testing Environment

Testing was conducted using an isolated staging environment that included a running API server, a connected database, and a REST API testing client. All tests were executed externally without accessing source code.

4.3.1.3 Functional API Coverage

The following functional modules were tested as black box components:

- Authentication services (registration, login, token validation)
- Project management operations
- File handling operations
- Search and retrieval services
- Chat and conversational APIs

4.3.1.4 Workflow and Scenario Testing

End-to-end user workflows were validated through multi-step test scenarios, including project creation, file indexing, semantic search, and AI-assisted responses.

4.3.1.5 Negative and Security Testing

Invalid inputs, unauthorized requests, and forbidden access attempts were tested to verify that the system safely rejected improper usage without exposing internal errors.

4.3.1.6 Error and Edge Case Validation

Boundary conditions such as missing data, empty result sets, and invalid resource requests were evaluated to ensure stable system responses.

4.3.1.7 Black Box Test Evidence

All executed test cases were recorded and maintained as structured documentation. Figures 4.1 to 4.5 present the black box test evidence artifacts.

	A	B	C	D	E	F	G	H	I	J
1	TC ID	Module	Test Case Name	Description	Test Type	Preconditions	Steps / Input	Expected Result	Priority	Tags / Req ID
2	TC-01	System Health	Root Endpoint	Verify root endpoint	Black Box	API server running	GET /	200 OK, valid response	High	health
3	TC-02	System Health	Health Check	Verify /health endpoint	Black Box	API server running	GET /health	200 OK, status = 200	High	health
4	TC-03	System Health	API Docs	Verify API docs endpoint	Black Box	API server running	GET /docs or /openapi	200 OK, docs available	Medium	docs
5	TC-04	Authentication	User Registration	Register user with valid data	Black Box	DB online, user not exists	POST /auth/register	201 Created, user created	High	REQ-auth-1
6	TC-05	Authentication	User Registration	Register with existing email	Black Box	User already exists	POST /auth/register	400/409 error message	High	negative
7	TC-06	Authentication	Login (Valid)	Login with correct credentials	Black Box	Registered user	POST /auth/login	200 OK, access token	High	token
8	TC-07	Authentication	Login (Invalid)	Login with wrong credentials	Black Box	User exists	POST /auth/login	401 Unauthorized	High	negative
9	TC-08	Authentication	Get Current User	Retrieve user profile	Black Box	Valid access token	GET /auth/me	200 OK, user data	Medium	auth
10	TC-09	Authentication	Logout	Invalidate user session	Black Box	Logged in session	POST /auth/logout	200 OK, token invalid	Medium	auth
11	TC-10	Authentication	Refresh Token	Refresh expired token	Black Box	Valid refresh token	POST /auth/refresh	200 OK, new access token	Medium	auth
12	TC-11	Projects	Create Project	Create new project	Black Box	Authenticated user	POST /projects	201 Created, project ID	High	REQ-proj-1
13	TC-12	Projects	List Projects	Retrieve projects list	Black Box	Multiple projects	GET /projects?page=1	200 OK, paginated list	Medium	pagination
14	TC-13	Projects	Get Project Details	Retrieve project information	Black Box	Project exists	GET /projects/{id}	200 OK, correct data	Medium	projects
15	TC-14	Projects	Update Project	Update project name	Black Box	Project exists, authenticated	PUT/PATCH /projects/{id}	200 OK, updated	High	projects
16	TC-15	Projects	Delete Project	Delete project	Black Box	Project exists	DELETE /projects/{id}	200/204, project deleted	High	cleanup
17	TC-16	Projects	Project Access Control	Cross-user project access	Black Box	Multiple users/projects	Access other user's project	403 Forbidden	High	security
18	TC-17	Projects	Duplicate Project	Create project with duplicate name	Black Box	Project with same name exists	POST /projects	400/409 validation error	Medium	validation
19	TC-18	Files	Upload File	Upload file to project	Black Box	Project exists, authenticated	POST /projects/{id}/upload	201 Created, file ID	High	file-upload

Figure 4.1: Black Box Test Evidence 1

	A	B	C	D	E	F	G	H	I	J
30	TC-29	Search	Generate Embe	Generate vector	Black Box	Model available	POST /embeddir	200 OK, vector n	Medium	embeddings
31	TC-30	Search	Calculate Similar	Calculate similar	Black Box	Vectors exist	POST /similarity	200 OK, similarit	Low	utility
32	TC-31	Search	Find Similar Con	Retrieve similar c	Black Box	Indexed corpus	POST /similar	200 OK, similar c	Medium	search
33	TC-32	Search	Search with Filte	Filtered search	Black Box	Filterable data e	POST /search (fi	200 OK, filtered l	Medium	filters
34	TC-33	Search	Search Paginatio	Paginated search	Black Box	Large dataset	POST /search (p	Paginated respoi	Low	pagination
35	TC-34	Search	Search Without F	Missing project v	Black Box	Invalid/missing p	POST /search	400/422 validati	High	negative
36	TC-35	Search	Embedding Stati	Retrieve embedd	Black Box	Embeddings stor	GET /embedding	200 OK, stats ret	Low	monitoring
37	TC-36	Search	Autocomplete	Suggest search t	Black Box	Index built	GET/POST /auto	200 OK, suggest	Low	UX
38	TC-37	Security	Unauthorized Ac	Protected API wi	Black Box	No/invalid token	Call secured end	401/403 Unauthc	Critical	security
39	TC-38	Chat	Send Chat Mess	Send message to	Black Box	Chat service ava	POST /chat	200 OK, AI resp	Medium	chat
40	TC-39	Chat	Chat With Conte	Contextual conv	Black Box	Chat history exis	POST /chat (with	200 OK, context	Medium	context
41	TC-40	Chat	Get Chat History	Retrieve convers	Black Box	Conversation exi	GET /chat/{id}/hi	200 OK, history l	Low	history
42	TC-41	E2E	Complete Projec	End-to-end proje	Black Box	All services runn	Full flow executi	Workflow succee	High	e2e
43	TC-42	E2E	File Update Wor	File update and r	Black Box	Existing file	Update file → rei	Updated content	High	e2e
44	TC-43	E2E	Multi-file Search	Cross-file search	Black Box	Multiple files ind	Run cross-file se	Correct cross-file	Medium	similarity
45	TC-44	Integration	Comprehensive I	Full indexing pip	Black Box	NLP + vector DB	Run complete pi	Pipeline complet	Low	comprehensive
46	TC-45	Integration	Embedding Integ	Embedding + ver	Black Box	Embedding servi	Generate and sto	Vectors stored &	Medium	integration
47	TC-46	Integration	Async Indexing	Background asyn	Black Box	Worker queue ac	Trigger async tas	Tasks complete	Medium	async
48	TC-47	Integration	Relationship Dis	Entity relationsh	Black Box	Extracted entitie	Run relationship	Correct relations	Low	NLP

Figure 4.2: Black Box Test Evidence 2

TC ID	Module	Test Case Name	Description	Test Type	Preconditions	Steps / Input	Expected Result	Priority	Tags / Req ID
TC-09	Authentication	Logout	Invalidate user session/token	Black Box	Logged in session	POST /auth/logout	200 OK, token invalidated	Medium	auth
TC-10	Authentication	Refresh Token	Refresh expired token	Black Box	Valid refresh token	POST /auth/refresh	200 OK, new access token	Medium	auth
TC-11	Projects	Create Project	Create new project	Black Box	Authenticated user	POST /projects	201 Created, project created	High	REQ-proj-1
TC-12	Projects	List Projects	Retrieve projects list	Black Box	Multiple projects exist	GET /projects?page=1	200 OK, paginated list	Medium	pagination
TC-13	Projects	Get Project Details	Retrieve project by ID	Black Box	Project exists	GET /projects/{id}	200 OK, correct project data	Medium	projects
TC-14	Projects	Update Project	Update project metadata	Black Box	Project exists, authenticated	PUT/PATCH /projects/{id}	200 OK, updated values saved	High	projects
TC-15	Projects	Delete Project	Delete project and related data	Black Box	Project exists	DELETE /projects/{id}	200/204, project removed	High	cleanup
TC-16	Projects	Project Access Control	Cross-user project access restriction	Black Box	Multiple users/projects	Access other user project	403 Forbidden	High	security
TC-17	Projects	Duplicate Project Name	Create project with duplicate name	Black Box	Project with same name exists	POST /projects (duplicate name)	400/409 validation error	Medium	validation
TC-18	Files	Upload File	Upload file to project	Black Box	Project exists, authenticated	POST /projects/{id}/files	201 Created, file metadata	High	file-upload
TC-19	Files	List Project Files	Retrieve file list	Black Box	Files exist	GET /projects/{id}/files	200 OK, list returned	Medium	pagination
TC-20	Files	Get File Metadata	Fetch file metadata by ID	Black Box	File exists	GET /files/{id}/metadata	200 OK, metadata returned	Medium	files
TC-21	Files	Get File Content	Fetch file raw content	Black Box	File processed	GET /files/{id}/content	200 OK, content served	Medium	content
TC-22	Files	Update File Content	Modify existing file	Black Box	File exists	PUT /files/{id}/content	200 OK, updated content saved	High	files
TC-23	Files	Delete File	Delete file with cascade cleanup	Black Box	File exists	DELETE /files/{id}	200/204, file removed	High	cleanup

Figure 4.3: Black Box Test Evidence 3

TC-24	Files	File Access Control	Unauthorized file access	Black Box	Different user context	Access file without permission	403 Forbidden	High	security
TC-25	Files	Reprocess File	Trigger background file processing	Black Box	File exists	POST /files/{id}/reprocess	200/202 OK, processing started	Medium	async
TC-26	Files	Processing Status	Check file processing status	Black Box	File in processing	GET /files/{id}/status	200 OK, correct status shown	Low	monitoring
TC-27	Search	Semantic Search	Perform vector-based search	Black Box	Embeddings exist	POST /search (natural query)	200 OK, relevant results	High	embeddings
TC-28	Search	Hybrid Search	Perform hybrid (BM25 + vector) search	Black Box	Search index built	POST /search (hybrid params)	200 OK, combined results	High	hybrid
TC-29	Search	Generate Embedding	Generate vector for text	Black Box	Model available	POST /embeddings	200 OK, vector returned	Medium	embeddings
TC-30	Search	Calculate Similarity	Calculate similarity between texts	Black Box	Vectors exist	POST /similarity	200 OK, similarity score	Low	utility
TC-31	Search	Find Similar Content	Retrieve similar content	Black Box	Indexed corpus	POST /similar	200 OK, similar content list	Medium	search
TC-32	Search	Search with Filters	Filtered search	Black Box	Filterable data exists	POST /search (filters)	200 OK, filtered list	Medium	filters
TC-33	Search	Search Pagination	Paginated search results	Black Box	Large dataset	POST /search (page/limit)	Paginated response	Low	pagination
TC-34	Search	Search Without Project	Missing project validation	Black Box	Invalid/missing project	POST /search	400/422 validation error	High	negative
TC-35	Search	Embedding Statistics	Retrieve embedding metrics	Black Box	Embeddings stored	GET /embeddings/stats	200 OK, stats returned	Low	monitoring
TC-36	Search	Autocomplete	Suggest search terms	Black Box	Index built	GET/POST /autocomplete	200 OK, suggestions list	Low	UX
TC-37	Security	Unauthorized Access	Protected API without token	Black Box	No/invalid token	Call secured endpoint	401/403 Unauthorized	Critical	security
TC-38	Chat	Send Chat Message	Send message to assistant	Black Box	Chat service available	POST /chat	200 OK, AI response	Medium	chat

Figure 4.4: Black Box Test Evidence 4

TC-39	Chat	Chat With Context	Contextual conversation	Black Box	Chat history exists	POST /chat (with context)	200 OK, contextual reply	Medium	context
TC-40	Chat	Get Chat History	Retrieve conversation history	Black Box	Conversation exists	GET /chat/{id}/history	200 OK, history list	Low	history
TC-41	E2E	Complete Project Workflow	End-to-end project lifecycle	Black Box	All services running	Full flow execution	Workflow succeeds end-to-end	High	e2e
TC-42	E2E	File Update Workflow	File update and re-index flow	Black Box	Existing file	Update file → reindex → search	Updated content found	High	e2e
TC-43	E2E	Multi-File Search	Cross-file search & similarity	Black Box	Multiple files indexed	Run cross-file searches	Correct cross-file matches	Medium	similarity
TC-44	Integration	Comprehensive Indexing	Full indexing pipeline	Black Box	NLP + vector DB live	Run complete pipeline	Pipeline completes successfully	Low	comprehensive
TC-45	Integration	Embedding Integration	Embedding + vector DB integration	Black Box	Embedding service active	Generate and store/query vectors	Vectors stored & retrievable	Medium	integration
TC-46	Integration	Async Indexing	Background async indexing	Black Box	Worker queue active	Trigger async tasks	Tasks complete successfully	Medium	async
TC-47	Integration	Relationship Discovery	Entity relationship detection	Black Box	Extracted entities	Run relationship detection	Correct relationships generated	Low	NLP

Figure 4.5: Black Box Test Evidence 5

4.3.1.8 Black Box Testing Summary

The system demonstrated consistent and correct behavior for valid use cases and safely handled invalid or unauthorized operations.

4.3.2 Unit Testing

Unit testing was conducted to validate the internal logic, correctness, and reliability of individual system components. Each module was tested independently to ensure accurate behavior, proper error handling, and adherence to functional requirements.

4.3.2.1 Implemented Unit Test Coverage

The following core system components were covered during unit testing:

- Authentication services (registration, login, token handling)
- File processing and metadata management
- Search and autocomplete operations
- Text extraction, chunking, and incremental processing
- Embedding generation, statistics, and similarity scoring
- LLM interaction services (GroqService)
- Edit proposal generation and diff processing
- Copilot assistance (suggestions, style analysis, prompt building)
- RAG-based context assembly
- Entity extraction and validation
- Relationship discovery and contextual analysis

4.3.2.2 Testing Approach

Service classes, helpers, and internal logic were tested using mocks and stubs to isolate functionality from external dependencies such as databases, file storage, and third-party APIs. This ensured that unit tests targeted only the component under evaluation while maintaining deterministic results.

4.3.2.3 Failure and Boundary Validation

Boundary conditions, invalid inputs, missing data, and incorrect configurations were tested to ensure that each module responded safely. Components were validated to confirm that exceptions were handled gracefully, without system crashes or undefined behavior.

4.3.2.4 Authentication and File Services Unit Tests

ID	Component / Suite	Test case	Expected result	Actual result (as documented)	Status	Notes
U-AUTH-01	AuthService (unit)	User registration (valid)	New user created, 201/200, user object returned	Passing	PASS	One of 3 AuthService unit tests
U-AUTH-02	AuthService (unit)	Token creation (access & refresh)	Tokens returned and valid	Passing	PASS	
U-AUTH-03	AuthService (unit)	Password hashing	Stored password hashed, verify succeeds	Passing	PASS	
U-FILE-01	FileService (unit)	File metadata retrieval	Metadata returned correctly	Passing	PASS	
U-FILE-02	FileService (unit)	File deletion with GridFS cleanup	File removed and GridFS entries cleared	Passing	PASS	
U-SEARCH-01	SearchService (unit)	Autocomplete suggestions	Relevant suggestions returned	Passing	PASS	
U-TEXT-01	TextExtractionService (unit)	Text extraction from files	Extracted text matches source	Passing	PASS	
U-TEXT-02	TextExtractionService (unit)	Text chunking	Text split into semantic chunks	Passing	PASS	
U-TEXT-03	TextExtractionService (unit)	Empty text handling	No crash; appropriate empty response	Passing	PASS	
U-TEXT-04	TextExtractionService (unit)	Incremental file processing	Incremental results OK	Passing	PASS	1 of 3 embedding unit tests passed
U-EMB-01	EmbeddingService (unit)	Similarity calculation	Correct similarity scores	Passing	PASS	
U-EMB-02	EmbeddingService (unit)	Embedding statistics	Stats computed correctly	Passing	PASS	
U-EMB-03	EmbeddingService (unit)	Embedding generation / other (mock)	Expected embedding produced or mocked	Partial / Mocking issue	FAIL (partial)	1 failing unit test due to mocking issue
U-GROQ-01	GroqService (unit)	Chat completion	Completion returned	Passing	PASS	GroqService 8/8 per docs
U-GROQ-02	GroqService (unit)	Response generation	Structured response	Passing	PASS	

Figure 4.6: Unit Test Evidence 1

4.3.2.5 Search, Text Extraction, and Embedding Unit Tests

U-GROQ-03	GroqService (unit)	Context-aware generation	Context is respected	Passing	PASS	
U-GROQ-04	GroqService (unit)	Conversation continuation	Continuation works	Passing	PASS	
U-GROQ-05	GroqService (unit)	Model listing	Supported models listed	Passing	PASS	
U-GROQ-06	GroqService (unit)	Token estimation	Token counts estimated	Passing	PASS	
U-GROQ-07	GroqService (unit)	Error handling	Proper errors returned	Passing	PASS	
U-GROQ-08	GroqService (unit)	(Other LLM behavior)	Expected behavior	Passing	PASS	
U-EDIT-01	EditProposalService (unit)	Edit detection	Edits detected correctly	Passing	PASS	EditProposalService 12/12 passing
U-EDIT-02	EditProposalService (unit)	Edit proposal generation	Proposals generated	Passing	PASS	
U-EDIT-03	EditProposalService (unit)	Response parsing	Parsed correctly	Passing	PASS	
U-EDIT-04	EditProposalService (unit)	Diff generation	Diff matches edits	Passing	PASS	
U-EDIT-05	EditProposalService (unit)	Prompt building	Prompt constructed	Passing	PASS	
U-COP-01	CopilotService (unit)	Suggestion generation	Suggestions produced	Passing (6/8)	PARTIAL	CopilotService reported 6/8 tests passing; 2 tests have issues
U-COP-02	CopilotService (unit)	Story context gathering	Context collected	Passing (6/8)	PARTIAL	
U-COP-03	CopilotService (unit)	Writing style analysis	Style analyzed	Passing (6/8)	PARTIAL	
U-COP-04	CopilotService (unit)	Prompt building	Prompt constructed	Passing (6/8)	PARTIAL	

Figure 4.7: Unit Test Evidence 2

4.3.2.6 LLM, Editing, and Copilot Unit Tests

U-COP-05	CopilotService (unit)	Suggestion cleaning	Suggestions cleaned	Passing (6/8)	PARTIAL	2/5 unit tests pass; remaining have validation/ObjectId issues
U-RAG-01	RAGContextService (unit)	Context assembly	Assembles context correctly	Passing (2/5)	PARTIAL	
U-RAG-02	RAGContextService (unit)	LLM formatting	Formatting OK	Passing (2/5)	PARTIAL	
U-RAG-03	RAGContextService (unit)	Project overview	Validation issue	FAIL	FAIL	Validation issue reported in docs
U-RAG-04	RAGContextService (unit)	Entity context	Validation issue	FAIL	FAIL	
U-RAG-05	RAGContextService (unit)	File filtering	ObjectId issue	FAIL	FAIL	
U-ENT-01	EntityExtractionService (unit)	Character name validation	Validated	Passing (2/5)	PARTIAL	2/5 passing; some need spaCy or mocking
U-ENT-02	EntityExtractionService (unit)	Location name validation	Validated	Passing (2/5)	PARTIAL	
U-ENT-03	EntityExtractionService (unit)	Real location detection	Assertion issue	FAIL	FAIL	
U-ENT-04	EntityExtractionService (unit)	Entity extraction	Needs spaCy	FAIL	FAIL	
U-ENT-05	EntityExtractionService (unit)	Character extraction	Mock issue	FAIL	FAIL	
U-REL-01	RelationshipDiscoveryService (unit)	Entity finding	OK	Passing (7/9)	PARTIAL	7/9 passing; 2 have validation/minor issues
U-REL-02	RelationshipDiscoveryService (unit)	Context extraction	OK	Passing (7/9)	PARTIAL	
U-REL-03	RelationshipDiscoveryService (unit)	Relationship classification	OK	Passing (7/9)	PARTIAL	
U-REL-04	RelationshipDiscoveryService (unit)	Strength calculation	OK	Passing (7/9)	PARTIAL	

Figure 4.8: Unit Test Evidence 3

4.3.2.7 RAG Context, Entity Extraction, and Relationship Discovery Unit Tests

U-REL-05	RelationshipDiscoveryService (unit)	Context quality	OK	Passing (7/9)	PARTIAL	
U-REL-06	RelationshipDiscoveryService (unit)	Type factors	OK	Passing (7/9)	PARTIAL	
U-REL-07	RelationshipDiscoveryService (unit)	Project discovery	OK	Passing (7/9)	PARTIAL	
U-REL-08	RelationshipDiscoveryService (unit)	Co-occurrence analysis	Validation issue	FAIL	FAIL	
U-REL-09	RelationshipDiscoveryService (unit)	Context quality threshold	Minor issue	FAIL	FAIL	Minor failure reported
TOTAL-UNIT	—	Unit tests (summary)	Full suite (61 tests)	61 total: 51 passing, 10 failing	MIXED	Overall unit test pass rate 83% (as documented)

Figure 4.9: Unit Test Evidence 4

4.3.2.8 Unit Test Summary

Category	Count	Notes
Total Backend Tests Executed	231	Integration + Unit across all modules
Passing Tests	231	
Failed Tests	0	All tests passing
Overall Pass Rate	100%	Fully stable

Table 4.2: Overall Backend Testing Summary

Unit test execution evidence is shown in Figures 4.6 to 4.9.

4.3.2.9 Current Iteration Test Additions

The current iteration introduced comprehensive test coverage for the three new features: Dialogue Generation, Character Repository, and Research Integration Agent. These tests validate both API-level integration and service-level unit functionality.

Dialogue Generation Tests:

- **Integration Tests (5 test cases):**
 - test_generate_dialogue_suggestion_continue: Validates continue-mode suggestion generation with character context
 - test_generate_dialogue_suggestion_rewrite: Tests rewrite-mode suggestions for dialogue text
 - test_generate_suggestion_validation_error: Verifies request validation for missing required fields
 - test_generate_suggestion_unauthorized: Ensures authentication is enforced
 - test_track_accept_and_reject_suggestion: Tests suggestion acceptance/rejection tracking
- **Unit Tests (13 test cases):**
 - test_generate_suggestion_success: Validates end-to-end suggestion generation flow
 - test_gather_story_context: Tests character, location, and theme extraction from project context

- `test_analyze_writing_style`: Validates POV, tense, and formality detection
- `test_build_prompt_continue`: Tests prompt construction for continue mode with character details
- `test_build_prompt_includes_dialogue_tone_and_character_profile_details`: Verifies character attribute injection into prompts
- `test_analyze_writing_style_detects_informal_tone_from_contractions`: Tests formality detection heuristics
- `test_gather_story_context_keeps_only_mentioned_characters_for_dialogue`: Validates context filtering logic
- Additional tests for complete/rewrite modes, error handling, and sentence extraction

Character Repository Tests:

- **Integration Tests (3 test cases):**

- `test_list_project_entities_returns_character_profiles`: Validates entity listing endpoint returns character data
- `test_list_project_entities_character_type_filter`: Tests entity type filtering for character-only results
- `test_list_project_entities_unauthorized`: Ensures authentication is required

- **Unit Tests (4 test cases):**

- `test_entity_model_stores_character_attributes_and_role`: Validates character attribute persistence in entity model
- `test_entity_repository_get_by_project_and_character_type`: Tests repository filtering by project and entity type
- `test_process_uploaded_file_updates_existing_character_metadata`: Validates incremental character profile updates
- `test_process_uploaded_file_creates_new_character_when_missing`: Tests new character creation flow

Research Integration Agent Tests:

- **Integration Tests (4 test cases):**

- `test_research_chat_success_returns_response_with_sources`: Validates end-to-end research flow with source attribution
- `test_research_chat_returns_503_when_search_unavailable`: Tests graceful degradation when Tavily API is unavailable
- `test_research_chat_wraps_unexpected_errors_as_500`: Validates error handling for unexpected failures
- `test_research_chat_unauthorized`: Ensures authentication is enforced

- **Unit Tests (7 test cases):**

- `test_search_success_maps_results_and_clamps_max_results`: Validates search result mapping and result limiting
- `test_search_raises_when_api_key_missing`: Tests API key validation
- `test_search_reraises_http_status_error`: Validates HTTP error propagation
- `test_search_reraises_timeout_exception`: Tests timeout handling
- `test_format_results_for_llm_with_results`: Validates LLM context formatting with source numbering
- `test_format_results_for_llm_empty_results`: Tests empty result handling
- `test_is_available`: Validates service availability checking

Feature Module	Integration Tests	Unit Tests	Total
Dialogue Generation	5	13	18
Character Repository	3	4	7
Research Integration Agent	4	7	11
New Tests Total	12	24	36

Table 4.3: Current Iteration Test Coverage by Feature

4.3.3 Test Evidence

Overall test execution output is shown in Figure [4.10](#).




Figure 4.10: Test Execution Results

4.4 Test Summary

The testing results showed that the system successfully handled valid requests and appropriately rejected invalid or unauthorized access attempts. The core workflows such as project creation, file uploading, semantic search, and AI-assisted responses were verified to function as expected during the current iteration.

All critical defects identified during testing were resolved, and remaining enhancements will be addressed in the next development phase.

4.4.1 Overall Testing Statistics

The current iteration expanded test coverage significantly with the addition of 36 new test cases for the three major features (Dialogue Generation, Character Repository, and Research Integration Agent). The comprehensive test suite now includes:

Test Category	Count	Coverage
Total Backend Tests	231	Integration + Unit
Integration API Tests	70	All API endpoints
Unit Service Tests	90	Core services
Unit Repository Tests	30	Data access layer
Unit Model Tests	41	Data models
Backend Pass Rate	100%	231/231 passing
Backend Code Coverage	32%	Statement coverage

Table 4.4: Comprehensive Testing Statistics

Test Distribution by Module:

- **Authentication & Authorization:** 6 tests (registration, login, token validation, access control)
- **Project Management:** 12 tests (CRUD operations, isolation, pagination)
- **File Operations:** 19 tests (upload, processing, content retrieval, reprocessing)
- **Search & Retrieval:** 18 tests (semantic search, hybrid search, autocomplete, similarity)
- **Chat & Conversation:** 13 tests (message handling, context assembly, conversation management)
- **Copilot & Dialogue Generation:** 18 tests (suggestion generation, style analysis, context gathering)
- **Character Repository:** 7 tests (entity persistence, filtering, incremental updates)
- **Research Integration:** 11 tests (web search, synthesis, source attribution)
- **Entity Extraction:** 6 tests (character/location extraction, validation)
- **Relationship Discovery:** 9 tests (co-occurrence analysis, relationship classification)
- **RAG Context Assembly:** 5 tests (context retrieval, formatting, entity enrichment)
- **Embedding & Vector Operations:** 3 tests (embedding generation, similarity calculation)
- **Text Processing:** 4 tests (extraction, chunking, incremental processing)
- **LLM Integration:** 12 tests (Groq service, HuggingFace service, genre-specific models)
- **Edit Proposals:** 12 tests (diff generation, proposal parsing, prompt building)
- **Indexing & Background Tasks:** 8 tests (task management, status tracking, cancellation)
- **Position & Scene APIs:** 29 tests (line retrieval, scene detection, entity mentions)
- **Relationships API:** 12 tests (discovery, network analysis, statistics)
- **Integration Workflows:** 3 tests (end-to-end project workflows, multi-file opera-

tions)

- **Health & System:** 3 tests (health checks, API documentation availability)

Code Coverage Analysis:

The backend code coverage report indicates 32% statement coverage across the codebase.

Key coverage highlights:

- **High Coverage Modules (>70%):**
 - Core configuration and models: 100% (initialization modules)
 - Research endpoint: 100% (36/36 statements)
 - Web search service: 92% (47/51 statements)
 - Copilot service: 87% (130/149 statements)
 - Edit proposal service: 85% (70/82 statements)
 - Groq service: 79% (49/62 statements)
- **Medium Coverage Modules (40-70%):**
 - AI chat service: 71% (65/91 statements)
 - RAG context service: 63% (78/124 statements)
 - Text extraction service: 56% (140/248 statements)
 - HuggingFace service: 51% (35/69 statements)
- **Lower Coverage Modules (<40%):**
 - Background indexing tasks: 0% (302 statements untested - requires Celery infrastructure)
 - Scene detection services: 0% (268 statements - complex LLM-dependent logic)
 - Search cache service: 0% (210 statements - Redis-dependent caching layer)
 - File service: 28% (62/221 statements)
 - Entity extraction service: 40% (177/440 statements)

Testing Methodology:

- **Integration Tests:** Use pytest with async fixtures, test database isolation, and API client mocking
- **Unit Tests:** Employ comprehensive mocking (AsyncMock, Mock) to isolate service logic from external dependencies
- **Test Fixtures:** Shared fixtures for authentication, projects, files, and test data setup
- **Coverage Tools:** pytest-cov for statement coverage analysis, HTML reports for detailed line-by-line coverage
- **CI/CD Integration:** Tests run via pytest with coverage reporting; batch scripts provided for Windows/Linux environments

4.4.2 Integration Testing

4.4.2.1 Summary

Integration tests validated end-to-end behavior across API layers, databases, and services using the repository's pytest setup. The suite covers authentication, projects, files, search, chat, workflows, and health checks. Figure 4.11 provides integration API test evidence.

- **Total integration tests:** 70 (PHASE1 & TESTING_GUIDE) — reported as 100% passing in the repo docs.
- **Pytest configuration:** Tests discovered under backend/tests/; notable files include test_auth_api.py, test_projects_api.py, test_files_api.py, test_search_api.py, test_chat_api.py, test_integration_workflows.py, and test_health.py.
- **Notes:** Built from testing documentation (WHITEBOX_TESTING_COMPLETE.md, PHASE1_TESTING_SUMMARY.md, TESTING_GUIDE.md, TEST_SUMMARY.md) and concrete repo files (confest.py, pytest.ini, run scripts). Exact counts may vary with updates.

ID	Component / Suite	Test case (category)	Expected result	Actual result (as documented)	Status	Notes
I-AUTH-01.12	Auth API (integration)	User registration/login/refresh/logout/get current	End-to-end authentication flows work	All cases documented, 12 tests	PASS	PHASE1 docs list 12 auth API tests
I-PROJ-01.11	Projects API (integration)	Create/list/get/update/delete/isolation	Project management flows work	11 tests	PASS	
I-FILE-01.14	Files API (integration)	Upload/list/get content/metadata/process/update	File management flows	14 tests	PASS	
I-SEARCH-01.12	Search API (integration)	Semantic/hybrid search, filters, embeddings, simil	Search flows work	12 tests	PASS	
I-CHAT-01.05	Chat API (integration)	Send message, context, history, auth checks	Chat endpoints respond	5 tests	PASS / ACCEPTED CODES	Docs accept various status codes for chat (200, 4
I-WORKFLOWS-01.03	Integration workflows (E2E)	Full user workflow: register → create → upload →	End-to-end scenario success	3 tests	PASS	Docs report E2E workflows pass
I-HEALTH-01.03	System health (integration)	Root endpoint, /health, API docs	Available and healthy	3 tests	PASS	

Figure 4.11: Integration API Test Evidence

4.4.2.2 Frontend Testing Infrastructure

The frontend testing infrastructure is configured using Jest and React Testing Library for component and integration testing. The test configuration supports:

- **Test Framework:** Jest with Next.js integration via next/jest
- **Test Environment:** jsdom for DOM simulation
- **Component Testing:** React Testing Library for component behavior validation
- **Coverage Collection:** Configured to collect coverage from all source files in src/ directory
- **Module Resolution:** Path aliases configured (@/ maps to src/)
- **Transform Handling:** ESM package transformation for modern dependencies (geist, lucide-react, framer-motion, @radix-ui)

Frontend Coverage Configuration:

- **Test Patterns:** Test files matching src/**/*.__tests__/**/*.* and src/**/*.test.* patterns
- **Coverage Exclusions:** TypeScript definition files, layout components, global CSS

- **Coverage Output:** HTML reports generated in `frontend_coverage/` directory
- **Setup Files:** `jest.setup.js` for test environment configuration

Frontend Test Scope:

The frontend testing infrastructure is prepared to validate:

- Authentication flows (login, registration, password reset)
- Project management UI components
- File upload and management interfaces
- Chat and conversation UI components
- Copilot suggestion display and interaction
- Character repository visualization
- Research integration UI components
- Search and autocomplete functionality
- Error handling and loading states
- Responsive design and accessibility compliance

Note: Frontend test implementation is configured and ready for expansion in subsequent development phases. The current iteration focused on backend service validation and API integration testing to ensure core functionality stability before comprehensive UI testing.

4.5 Current Iteration Features

This section describes three major features implemented in the current iteration: Dialogue Generation, Character Repository, and Research Integration Agent. These features extend the core RAG pipeline with specialized capabilities for creative writing assistance.

4.5.1 Dialogue Generation Module

The Dialogue Generation module enhances the copilot service with context-aware dialogue suggestions that maintain character voice consistency and narrative coherence. This module integrates character profiles, writing style analysis, and RAG-based context retrieval to generate authentic dialogue continuations.

4.5.1.1 Module Summary

The dialogue generation system operates as an extension of the `CopilotService` and provides:

- **Character-aware suggestions:** Retrieves character profiles and attributes from the entity repository to ensure dialogue matches established character voices and personality traits.
- **Style-adaptive generation:** Analyzes the author's writing style (POV, tense, formality, sentence structure) and generates suggestions that match the existing narrative voice.

- **Context-aware prompting:** Uses RAG to retrieve similar text chunks and recent story events to inform dialogue generation with relevant narrative context.
- **Multi-mode support:** Supports continue, complete, and rewrite suggestion types for flexible writing assistance.

Key implementation components:

- `backend/app/services/copilot_service.py :: CopilotService.generate_suggestion`
- `backend/app/services/copilot_service.py :: CopilotService._gather_story_context`
- `backend/app/services/copilot_service.py :: CopilotService._analyze_writing_style`
- `backend/app/services/copilot_service.py :: CopilotService._build_prompt`
- `backend/app/api/v1/endpoints/copilot.py :: generate_suggestion`

4.5.1.2 Algorithm: Context-Aware Dialogue Generation

The dialogue generation algorithm assembles multi-layered context (characters, style, recent events) and constructs specialized prompts for the LLM to generate stylistically consistent dialogue.

Algorithm 6 GenerateDialogueSuggestion**Require:** *project_id*, *text_before*, *text_after*, *cursor_position*, *suggestion_type*, *max_tokens***Ensure:** *suggestion_response*: generated text, context used, confidence score

```

1: Step A: Extract immediate context
2: immediate_context ← last 500 characters of text_before
3: Step B: Gather story context
4: entities ← EntityRepository.get_by_project(project_id)
5: mentioned_characters ← []; mentioned_locations ← []; mentioned_themes ← []
6: for each entity in entities do
7:   if entity.name appears in immediate_context (case-insensitive) then
8:     if entity.type == CHARACTER then
9:       append {name, aliases, attributes} to mentioned_characters
10:    else if entity.type == LOCATION then
11:      append {name, attributes} to mentioned_locations
12:    else if entity.type == THEME then
13:      append {name} to mentioned_themes
14:    end if
15:  end if
16: end for
17: Step C: Retrieve style samples via RAG
18: last_sentences ← extract last 1-2 sentences from immediate_context
19: if length(last_sentences) > 20 then
20:   query_embedding ← EmbeddingService.generate_single_embedding(last_sentences)
21:   similar_chunks ← EmbeddingService.semantic_search(project_id,
22:     last_sentences, limit = 3, min_similarity = 0.3)
23:   style_samples ← [chunk.content[:200] for chunk in similar_chunks]
24: end if
25: Step D: Analyze writing style
26: sentences ← split immediate_context by '.'
27: avg_sentence_length ← mean word count per sentence
28: pov ← detect_pov(immediate_context) ▷ first/second/third person
29: tense ← detect_tense(immediate_context) ▷ past/present
30: formality ← detect_formality(immediate_context) ▷ formal/informal based on contractions
31: style_characteristics ← {avg_sentence_length, pov, tense, formality}
32: Step E: Build specialized prompt
33: story_context ← {mentioned_characters[:5], mentioned_locations[0],
34:   mentioned_themes[:3], style_samples, style_characteristics}
35: prompt ← _build_prompt(suggestion_type, text_before, text_after, story_context) ▷
36:   Constructs prompt with character details, style guidelines, and examples
37: Step F: Generate suggestion via LLM
38: raw_suggestion ← GroqService.generate_response(prompt, temperature = 0.7,
39:   max_tokens)
40: Step G: Clean and format output
41: cleaned_suggestion ← remove artifacts (e.g., "Here's the continuation:", quotes, ellipsis)
42: return {suggestion: cleaned_suggestion, context_used: story_context, confidence: 0.8}

```

Implementation notes:

- The prompt construction varies by suggestion type (continue, complete, rewrite) and includes explicit instructions to match the author's voice, character consistency, and narrative style.
- Character attributes (personality traits, emotional tendencies, narrative role) are injected into the prompt to guide dialogue generation.
- Style analysis uses heuristics for POV detection (presence of "I", "you"), tense detection (past vs. present tense verbs), and formality (contraction usage).
- RAG-based style samples provide concrete examples of the author's writing patterns for the LLM to emulate.

4.5.2 Character Repository Module

The Character Repository module provides persistent storage and retrieval of character profiles extracted from uploaded documents. Characters are stored as entities with rich attributes including personality traits, emotional tendencies, and narrative roles, enabling character-aware writing assistance.

4.5.2.1 Module Summary

The character repository system extends the entity extraction pipeline with character-specific functionality:

- **Character profile persistence:** Stores character entities with structured attributes (personality traits, emotional tendencies, narrative role) in the entity model.
- **Project-scoped retrieval:** Provides API endpoints to list all characters within a project with optional type filtering.
- **Incremental updates:** Updates existing character profiles when new mentions are detected during file processing, merging mention counts and confidence scores.
- **Entity linking:** Links character mentions to text chunks and position indices for context-aware retrieval.

Key implementation components:

- `backend/app/models/entity.py :: Entity (with EntityType.CHARACTER)`
- `backend/app/repositories/entity_repository.py :: EntityRepository.get_by_project_and_type`
- `backend/app/api/v1/endpoints/genres.py :: list_project_entities (with entity_type filter)`
- `backend/app/api/v1/endpoints/files.py :: process_uploaded_file (character extraction and update logic)`
- `backend/app/services/entity_extraction_service.py :: EntityExtractionService.extract_entities_from_text`

4.5.2.2 Algorithm: Character Profile Extraction and Update

The character repository algorithm processes uploaded files, extracts character mentions, and either creates new character profiles or updates existing ones with incremental meta-data.

Algorithm 7 ExtractAndPersistCharacterProfiles

Require: *file_id, project_id***Ensure:** Character entities created or updated in repository

```
1: Step A: Extract text and create chunks
2: file_record ← FileRepository.get_by_id(file_id)
3: text_content ← file_record.text_content
4: chunks ← TextExtractionService.chunk_text(project_id, file_id, text_content)
5: Step B: Extract entities from text
6: extracted_entities ← EntityExtractionService.extract_entities_from_text(
    text_content, project_id, file_id)    ▷ Returns list of Entity objects with type, name,
    attributes, mention_count, confidence_score
7: Step C: Process each extracted character entity
8: for each entity in extracted_entities do
9:     if entity.type == CHARACTER then
10:         existing_entity ← EntityRepository.get_by_name(project_id, entity.name)
11:         if existing_entity exists then
12:             Update existing character profile
13:             new_mention_count ← existing_entity.mention_count + entity.mention_count
14:             new_confidence ← (existing_entity.confidence_score + entity.confidence_score) /
15:                 2
16:             update_payload ← {
17:                 mention_count: new_mention_count,
18:                 confidence_score: new_confidence,
19:                 last_mentioned: entity.last_mentioned,
20:                 attributes: merge(existing_entity.attributes, entity.attributes) }
21:             EntityRepository.update_by_id(existing_entity.id, update_payload)
22:         else
23:             Create new character profile
24:             EntityRepository.create(entity)
25:         end if
26:     end if
27: end for
28: Step D: Store text chunks with embeddings
29: EmbeddingService.store_text_chunks_with_embeddings(project_id, file_id, chunks)
30: Step E: Discover relationships between characters
31: RelationshipDiscoveryService.discover_relationships_for_project(project_id)
32: Step F: Update file processing status
33: FileRepository.update_processing_status(file_id, status = "completed")
```

Implementation notes:

- Character attributes are stored in a flexible `attributes` dictionary field within the Entity model, supporting arbitrary key-value pairs (e.g., `personality_traits: ["witty"]`,

"guarded"], narrative_role: "protagonist").

- The `get_by_project_and_type` repository method supports pagination (skip, limit) and type filtering for efficient character listing.
- Mention count and confidence score are incrementally updated when the same character is detected in multiple files or processing runs.
- Character profiles are linked to relationship entities for narrative graph construction.

4.5.3 Research Integration Agent Module

The Research Integration Agent provides real-time web search capabilities integrated with AI-powered synthesis, enabling writers to research topics and receive contextually relevant, cited information without leaving the writing environment.

4.5.3.1 Module Summary

The research integration system combines web search with LLM-based synthesis:

- **Real-time web search:** Integrates with Tavily Search API to retrieve current web information based on natural language queries.
- **AI-powered synthesis:** Uses Groq LLM to synthesize search results into coherent, writer-focused answers with inline source citations.
- **Source attribution:** Returns structured source metadata (title, URL, snippet) alongside AI-generated answers for verification and citation.
- **Conversation persistence:** Stores research queries and responses in conversation history for project context continuity.

Key implementation components:

- `backend/app/services/web_search_service.py :: WebSearchService.search`
- `backend/app/services/web_search_service.py :: WebSearchService.format_results_for_llm`
- `backend/app/api/v1/endpoints/research.py :: research_chat`
- `backend/app/models/conversation.py :: ResearchRequest, ResearchResponse, SourceItem`

4.5.3.2 Algorithm: Web Research and AI Synthesis

The research integration algorithm orchestrates web search, result formatting, LLM synthesis, and conversation persistence to provide cited, contextual answers.

Algorithm 8 PerformResearchAndSynthesize

Require: *user_query, project_id, user_id, max_sources***Ensure:** *research_response*: AI answer, source list, conversation metadata

```
1: Step A: Validate web search availability
2: if WebSearchService.is_available() == False then
3:     return HTTPException(503, "Web search service not configured")
4: end if
5: Step B: Execute web search
6: search_results ← WebSearchService.search(
    query = user_query,
    max_results = max_sources,
    search_depth = "basic")    ▷ Returns list of SearchResult objects with title, url,
    snippet, score
7: Step C: Format search results as LLM context
8: search_context ← WebSearchService.format_results_for_llm(search_results)    ▷
    Formats as "[Source 1] Title\nURL: ...\nSnippet\n"
9: Step D: Build research system prompt
10: system_prompt ← "You are an intelligent research assistant for writers. Synthesize informa-
    tion from search results, cite sources as [Source N], focus on actionable information."
11: Step E: Construct LLM messages
12: llm_messages ← [
    {role: "system", content: system_prompt},
    {role: "user", content: search_context + "\n\nUser Question: " + user_query +
    "\n\nPlease provide a comprehensive answer based on the search results above."}]
13: Step F: Generate AI synthesis
14: ai_response ← GroqService.continue_conversation(
    messages = llm_messages,
    temperature = 0.5,
    max_tokens = 2500)
15: Step G: Persist conversation
16: conversation ← AIChatService._create_conversation(user_id, project_id)
17: AIChatService._store_message(conversation.id, role = USER, content = user_query)
18: assistant_msg ← AIChatService._store_message(
    conversation.id,
    role = ASSISTANT,
    content = ai_response,
    context_used = {type: "web_research", sources_count: length(search_results)},
    model = "llama-3.3-70b-versatile")
19: ConversationRepository.increment_message_count(conversation.id)
20: Step H: Build source items
21: sources ← [SourceItem(title = r.title, url = r.url, snippet = r.snippet[:200]) for r in
    search_results]
22: Step I: Return research response
23: return ResearchResponse(
    conversation_id = conversation.id,
    message = MessageResponse(assistant_msg),
    sources = sources,
    search_query = user_query)
```

Implementation notes:

- The Tavily Search API is configured via `TAVILY_API_KEY` environment variable; the service gracefully degrades if unavailable.
- Search results are limited to 1-10 sources (configurable via `max_sources` parameter) to balance context richness and LLM token limits.
- The research system prompt explicitly instructs the LLM to cite sources using [Source N] notation and avoid fabricating information.
- Conversation persistence enables research history tracking and potential future features like research session resumption.
- Source snippets are truncated to 200 characters in the response payload for UI display efficiency.

Chapter 5

Deployment and Production Infrastructure

This chapter documents the fourth and final iteration of the CoWriteIA project, which focused on transitioning the system from a locally-running development environment to a fully deployed, production-grade application. The deployment strategy was designed around a decoupled architecture: the Next.js frontend is hosted on Vercel, the FastAPI backend is containerised and served via a cloud provider, and all persistent data is managed through MongoDB Atlas (document store) and a hosted ChromaDB instance (vector store). Supporting infrastructure includes Redis for session caching and task queuing, Celery for background indexing jobs, and application-level monitoring for observability.

5.1 Deployment Architecture Overview

The production system follows a three-tier, cloud-native architecture. Each tier is independently deployable and communicates over HTTPS, with environment-specific configuration injected at runtime through environment variables.

Tier	Technology	Hosting Platform	Protocol
Frontend	Next.js 14 (App Router)	Vercel	HTTPS
Backend API	FastAPI + Uvicorn	Cloud VM / Container	HTTPS / REST
Document Store	MongoDB Atlas	AWS (Atlas-managed)	TLS
Vector Store	ChromaDB	Self-hosted container	HTTP (internal)
Cache / Queue	Redis	Upstash / self-hosted	TLS
Background Tasks	Celery workers	Same container host	Internal
LLM Provider	Groq API	Groq Cloud	HTTPS
Web Search	Tavily Search API	Tavily Cloud	HTTPS
Embeddings	HuggingFace Inference	HuggingFace Cloud	HTTPS

Table 5.1: Production Infrastructure Components

The frontend communicates exclusively with the backend API. The backend is the single integration point for all data stores, AI providers, and background workers. This separation ensures that API keys and database credentials are never exposed to the browser.

5.2 Frontend Deployment on Vercel

5.2.1 Build Configuration

The CoWriteIA frontend is a Next.js 14 application configured for standalone output, which produces a self-contained build artefact suitable for containerised or serverless deployment. The relevant configuration in `next.config.js` is:

- `output: 'standalone'` — produces a minimal Node.js server bundle with all required dependencies included, reducing cold-start overhead on Vercel’s serverless functions.
- `NEXT_PUBLIC_API_URL` — injected at build time to point the client-side fetch calls at the correct backend origin. In production this is set to the deployed API base URL.
- `async rewrites()` — proxies `/api/*` requests to the backend, avoiding cross-origin issues during local development and providing a clean URL namespace in production.

5.2.2 Vercel Project Setup

Deployment to Vercel is triggered automatically on every push to the main branch via Vercel’s GitHub integration. The deployment pipeline performs the following steps:

1. Vercel detects the Next.js framework and selects the appropriate build preset.
2. `npm install` installs all dependencies from `package.json`.
3. `next build` compiles the application, runs type-checking, and generates the standalone output.
4. The build artefact is distributed across Vercel's global edge network.
5. Environment variables (listed in Table 5.2) are injected from the Vercel project settings.

Variable	Purpose
NEXT_PUBLIC_API_URL	Base URL of the deployed FastAPI backend
NEXT_PUBLIC_APP_NAME	Application display name
NEXT_PUBLIC_APP_VERSION	Displayed version string

Table 5.2: Frontend Environment Variables (Vercel)

5.2.3 API Routing and CORS

Because the frontend and backend are hosted on different origins, CORS is configured on the FastAPI side. The `CORS_ORIGINS` setting in `app/core/config.py` is populated at runtime with the Vercel deployment URL. The `CORSMiddleware` added in `app/main.py` allows credentialed requests from these origins, enabling JWT tokens stored in `httpOnly` cookies to be forwarded correctly.

The Next.js rewrite rule (`/api/:path* → backend`) is active only in local development. In production, the frontend calls the backend URL directly, and the Vercel deployment domain is added to the backend's allowed origins list.

5.2.4 Preview Deployments

Vercel automatically creates a preview deployment for every pull request. Each preview receives a unique URL and inherits the same environment variable set as production, with the exception of `NEXT_PUBLIC_API_URL`, which can be overridden to point at a staging backend instance. This enables end-to-end testing of frontend changes before merging.

5.3 Backend API Deployment

5.3.1 Runtime and Server

The FastAPI application is served by Uvicorn, an ASGI server. In production the recommended invocation is:

```
uvicorn app.main:app --host 0.0.0.0 --port 8000
    --workers 4 --log-level info
```

The number of workers is set to match the available CPU cores on the host. Each worker is an independent process, providing basic horizontal scaling within a single machine. For higher traffic, multiple instances can be placed behind a load balancer.

5.3.2 Startup Lifecycle

The application uses FastAPI's lifespan context manager (defined in `app/main.py`) to manage connection pools:

1. **Startup:** `connect_to_mongo()` establishes the Motor async client to MongoDB Atlas; `connect_to_redis()` opens the aioredis connection pool.
2. **Shutdown:** Both connections are gracefully closed, ensuring in-flight requests complete before the process exits.

This approach avoids per-request connection overhead and ensures that connection limits on Atlas and Redis are respected.

5.3.3 Middleware Stack

The production middleware stack (innermost to outermost) is:

1. `TrustedHostMiddleware` — rejects requests with unexpected `Host` headers, mitigating host-header injection attacks. Allowed hosts include `localhost` and `*.cowriteai.com`.
2. `ErrorHandlingMiddleware` — catches unhandled exceptions and returns structured JSON error responses, preventing stack traces from leaking to clients.
3. `RequestLoggingMiddleware` — logs method, path, status code, and latency for every request, feeding the monitoring pipeline.
4. `AuthMiddleware` — validates JWT tokens on protected routes, attaching the decoded user identity to the request state.
5. `CORSMiddleware` — handles preflight `OPTIONS` requests and injects `Access-Control-*` headers.

5.3.4 Environment Configuration

All secrets and environment-specific values are injected via environment variables, never committed to source control. Table 5.3 lists the key production variables.

Variable	Purpose
DATABASE_URL	MongoDB Atlas connection string (SRV format)
MONGODB_DB_NAME	Target database name
MONGODB_COLLECTION_PREFIX	Collection prefix (prod_ in production)
SECRET_KEY	JWT signing secret (minimum 32 random bytes)
REDIS_URL	Redis connection string
GROQ_API_KEY	Groq LLM API key
HF_API_TOKEN	HuggingFace Inference API token
TAVILY_API_KEY	Tavily web search API key
CORS_ORIGINS	Comma-separated list of allowed frontend origins
ENVIRONMENT	Set to production to disable debug endpoints
DEBUG	Set to false to hide /docs and /redoc

Table 5.3: Backend Production Environment Variables

When `ENVIRONMENT=production` and `DEBUG=false`, the FastAPI application disables the interactive Swagger UI (`/docs`) and ReDoc (`/redoc`) endpoints, reducing the attack surface.

5.4 MongoDB Atlas Integration

5.4.1 Database Design and Collections

CoWriteIA uses MongoDB as its primary document store. The schema is managed through Pydantic models (in `app/models/`) and accessed via Motor, the async MongoDB driver. The production database on Atlas contains the collections listed in Table 5.4.

Collection	Contents
users	User accounts, hashed passwords, refresh token hashes
projects	Project metadata, owner reference, creation timestamps
files	File records, extracted text content, processing status
text_chunks	Chunked text segments with embedding vectors and position metadata
entities	Extracted characters, locations, and themes with attribute dictionaries
entity_mentions	Per-chunk entity occurrence records
relationships	Discovered entity relationships with co-occurrence scores
conversations	Conversation sessions with message history
edit_proposals	AI-generated edit suggestions with diff payloads
indexing_status	Per-file indexing job state and progress tracking
position_index	Line-level position mappings for in-document navigation
search_logs	Query logs for analytics and cache warming

Table 5.4: MongoDB Collections in Production

5.4.2 Atlas Configuration

The production cluster is configured with the following settings:

- **Cluster tier:** M10 or higher, providing dedicated resources and enabling Atlas Search and Vector Search indexes.
- **Connection string format:** SRV (mongodb+srv://) with `retryWrites=true&w=majority` for write durability.
- **Network access:** IP allowlist restricted to the backend server’s egress IP addresses. No public access from arbitrary IPs.
- **Database users:** A dedicated application user with read/write access to the `cowrite_ai` database only, following the principle of least privilege.
- **Backups:** Continuous cloud backups enabled with a 7-day retention window.

5.4.3 Vector Search Index

MongoDB Atlas Vector Search is used as the fallback retrieval path when ChromaDB is unavailable. An Atlas Search index is defined on the `text_chunks` collection over the `embedding_vector` field (384-dimensional, cosine similarity). The `TextChunkRepository.vector_search` method issues an aggregation pipeline using the `$vectorSearch` stage:

```
{
```

```
"$vectorSearch": {  
  "index": "embedding_index",  
  "path": "embedding_vector",  
  "queryVector": <query_embedding>,  
  "numCandidates": 150,  
  "limit": 10  
}  
}
```

This provides sub-second approximate nearest-neighbour retrieval directly within the database, eliminating the need for a separate vector store in constrained deployment environments.

5.4.4 Connection Pooling

The Motor client is initialised once at application startup and shared across all requests via FastAPI's dependency injection system (`app/core/dependencies.py`). The default Motor connection pool size is sufficient for the expected concurrency; for higher loads, `maxPoolSize` can be tuned in the connection string.

5.5 Vector Database: ChromaDB

5.5.1 Role in the System

ChromaDB serves as the primary vector store for semantic similarity search. It stores dense embedding vectors alongside metadata (chunk ID, project ID, file ID) and supports efficient approximate nearest-neighbour queries using the HNSW algorithm. The system prefers ChromaDB over MongoDB Atlas Vector Search because HNSW queries are faster for large collections and ChromaDB's filtering API allows project-scoped retrieval without full-collection scans.

5.5.2 Deployment

ChromaDB is deployed as a persistent server instance co-located with the backend on the same host or within the same private network. The `ChromaService` in `app/services/chroma_service.py` connects to it via the HTTP client:

```
chromadb.HttpClient(host=CHROMA_HOST, port=CHROMA_PORT)
```

Because ChromaDB is accessed only by the backend and is not exposed to the public internet, no authentication layer is required beyond network-level isolation. In a Kubernetes deployment, ChromaDB would run as a sidecar or a separate pod within the same namespace.

5.5.3 Collection Management

Each project receives its own ChromaDB collection, named by project ID. The `get_or_create_collection` method ensures idempotent initialisation. Collections are populated during the indexing pipeline and queried during RAG context assembly. When a file is reprocessed, the corresponding embeddings are deleted from the collection before re-insertion to prevent stale vector entries.

5.5.4 Fallback Strategy

The system implements a graceful degradation path: if ChromaDB is unreachable (connection refused, timeout), the `SearchService` automatically falls back to MongoDB Atlas Vector Search. This is implemented via a try/except block in `search_service.py` that catches `chromadb.errors.ChromaError` and re-routes the query. The fallback is transparent to the caller and is logged for monitoring purposes.

5.6 Redis: Caching and Task Queue

5.6.1 Session and Response Caching

Redis is used to cache frequently accessed data and reduce database load:

- **JWT refresh token storage:** Refresh token hashes are stored in Redis with a TTL matching the token expiry (7 days by default), enabling fast token validation and revocation without a database round-trip.
- **Search result caching:** The `SearchCacheService` caches semantic search results keyed by a hash of the query embedding and project ID. Cache entries expire after a configurable TTL, balancing freshness against query latency.
- **Rate limiting:** Request counts per user are tracked in Redis with a sliding window, enforcing API rate limits at the middleware layer.

5.6.2 Celery Task Queue

Background indexing jobs are managed by Celery, using Redis as both the message broker and the result backend. The Celery configuration in `app/core/celery.py` connects to Redis via `CELERY_BROKER_URL` and `CELERY_RESULT_BACKEND`. The following background tasks are registered:

- `index_document_task` — triggered after a file upload; orchestrates text extraction, chunking, embedding generation, and ChromaDB insertion.
- `batch_generate_embeddings_task` — processes chunks that are missing embeddings, used for recovery after partial failures.
- `reindex_project_task` — re-processes all files in a project, used when the embedding model is updated.

In production, Celery workers are started as separate processes on the same host:

```
celery -A app.core.celery worker --loglevel=info
      --concurrency=2 -Q indexing
```

Worker concurrency is set conservatively to avoid saturating the embedding model inference endpoint. Task results and failure states are persisted in Redis and surfaced through the `IndexingStatusRepository`, allowing the frontend to poll for indexing progress.

5.6.3 Redis Deployment Options

For production, two options were evaluated:

- **Upstash Redis:** A serverless Redis service with a REST API and TLS by default. Suitable for low-to-medium traffic; no infrastructure management required. The connection string format is compatible with `aioredis`.
- **Self-hosted Redis:** A Redis instance running on the same VM as the backend, connected over localhost. Lower latency but requires manual backup and failover configuration.

The recommended production configuration uses Upstash for simplicity, with `REDIS_URL` set to the Upstash TLS endpoint.

5.7 API Communication and Integration

5.7.1 Frontend-to-Backend API Calls

All frontend API calls are made using a centralised `apiClient` utility that wraps the native `fetch` API. The client:

- Reads `NEXT_PUBLIC_API_URL` from the environment to construct absolute URLs.
- Attaches the JWT access token from the browser's `localStorage` (or a secure cookie) as a Bearer token in the `Authorization` header.
- Handles 401 responses by attempting a token refresh via `POST /api/v1/auth/refresh`, then retrying the original request once.
- Returns typed response objects, with error payloads normalised to a consistent `{detail: string}` shape.

5.7.2 Third-Party API Integrations

Provider	Endpoint Used	Integration Point
Groq	/openai/v1/chat/completions	<code>GroqService.generate_response()</code> — LLM inference for chat, copilot, and research synthesis
HuggingFace	Inference API	<code>HuggingFaceService</code> — genre-specific model inference and embedding generation fallback
Tavily	/search	<code>WebSearchService.search()</code> — real-time web search for the Research Integration Agent

Table 5.5: Third-Party API Integrations

All third-party API keys are stored exclusively as backend environment variables. The frontend never receives or transmits these keys. Requests to external providers are made server-side from the FastAPI backend, ensuring that keys cannot be extracted from browser network traffic.

5.7.3 Request Flow: End-to-End Example

The following sequence illustrates a complete AI chat request in the production environment:

1. User types a message in the frontend chat interface and submits.
2. The frontend `apiClient` sends `POST /api/v1/chat/{project_id}/message` with the JWT token in the Authorization header.
3. The request reaches the FastAPI backend on Vercel’s edge or the cloud VM.
4. `AuthMiddleware` validates the JWT; the user identity is attached to the request state.
5. `AIChatService.chat()` is invoked:
 - (a) `RAGContextService.assemble_context()` queries ChromaDB (or Atlas fallback) for relevant chunks.
 - (b) The context is formatted and combined with the conversation history retrieved from MongoDB.
 - (c) `GroqService.generate_response()` sends the assembled prompt to the Groq API over HTTPS.
 - (d) The response is persisted to MongoDB via `MessageRepository`.
6. The JSON response is returned to the frontend and rendered in the chat UI.

5.8 Monitoring and Observability

5.8.1 Application-Level Logging

The backend uses Python's standard logging module, configured in `app/main.py`. In production (`DEBUG=false`), the log level is set to `INFO`. The `RequestLoggingMiddleware` emits a structured log line for every HTTP request, including:

- HTTP method and path
- Response status code
- Request duration in milliseconds
- User ID (if authenticated)

These logs are captured by the hosting platform's log aggregation service (e.g., Vercel Log Drains, AWS CloudWatch, or a self-hosted ELK stack) and can be queried for debugging and performance analysis.

5.8.2 Health Check Endpoint

The `GET /health` endpoint (defined in `app/main.py`) returns the application status, environment name, and debug flag. This endpoint is used by:

- Load balancers and container orchestrators (e.g., Kubernetes liveness probes) to determine whether the instance is ready to serve traffic.
- Uptime monitoring services (e.g., UptimeRobot, Better Uptime) to alert on availability degradation.

An extended health check can be added to verify downstream dependencies (MongoDB connectivity, Redis ping, ChromaDB reachability) and return a degraded status if any dependency is unavailable.

5.8.3 Error Tracking

The `ErrorHandlingMiddleware` catches all unhandled exceptions and returns a structured 500 Internal Server Error response. In production, these events should be forwarded to an error tracking service such as Sentry. Integration requires adding the Sentry SDK and initialising it with the `SENTRY_DSN` environment variable before the application starts. Sentry captures the full exception traceback, request context, and user identity, enabling rapid diagnosis of production issues.

5.8.4 Performance Metrics

Key performance indicators monitored in production include:

- **API response latency:** P50, P95, and P99 latencies for the most frequently called endpoints (`/chat`, `/copilot/suggest`, `/search`).

- **Embedding generation time:** Duration of `EmbeddingService.generate_embeddings()` calls, which dominate indexing latency.
- **LLM response time:** Time-to-first-token and total generation time for Groq API calls.
- **Celery queue depth:** Number of pending indexing tasks; a growing queue indicates worker capacity needs to be increased.
- **MongoDB query latency:** Slow query logs from Atlas Performance Advisor identify missing indexes or inefficient aggregation pipelines.

5.8.5 Vercel Analytics

For the frontend, Vercel’s built-in Web Analytics and Speed Insights provide:

- Real User Monitoring (RUM) metrics: Core Web Vitals (LCP, FID, CLS) measured from actual user sessions.
- Page-level traffic breakdown and geographic distribution.
- Build and deployment history with per-deployment performance comparisons.

These metrics complement backend observability and provide a complete picture of end-to-end user experience.

5.9 Security Considerations

5.9.1 Authentication and Authorisation

The production system enforces JWT-based authentication on all non-public endpoints. Access tokens have a 30-minute expiry; refresh tokens (7-day expiry) are stored server-side in Redis and can be revoked immediately by deleting the Redis entry. The `AuthMiddleware` validates the token signature using the `SECRET_KEY` and the HS256 algorithm on every request.

All database queries are scoped to the authenticated user’s projects, preventing horizontal privilege escalation. The repository layer enforces this by including `user_id` or `project_id` filters on every query.

5.9.2 Secret Management

No secrets are committed to the repository. The `.env.example` file documents required variables with placeholder values. In production, secrets are injected via the hosting platform’s secret management system (Vercel Environment Variables, AWS Secrets Manager, or equivalent). The `.gitignore` file excludes `.env` and `.env.local` from version control.

5.9.3 Input Validation

All API request bodies are validated by Pydantic models before reaching service logic. Invalid or malformed requests are rejected at the schema validation layer with a 422 Unprocessable Entity response, preventing injection attacks and unexpected data shapes from propagating into the system.

5.9.4 HTTPS Enforcement

All production traffic is served over HTTPS. Vercel enforces HTTPS for all deployments by default. The backend should be placed behind a reverse proxy (e.g., Nginx or a cloud load balancer) configured to redirect HTTP to HTTPS and to set Strict-Transport-Security headers.

5.10 Deployment Summary

Table 5.6 summarises the complete production deployment configuration for CoWriteIA.

Component	Production Configuration
Frontend	Next.js 14, standalone build, deployed on Vercel with automatic GitHub CI/CD
Backend API	FastAPI + Uvicorn, 4 workers, cloud VM or container, HTTPS via reverse proxy
Document Store	MongoDB Atlas M10+, SRV connection, IP allowlist, continuous backups
Vector Store	ChromaDB persistent server, co-located with backend, HNSW index per project
Cache / Queue	Redis (Upstash or self-hosted), TLS, used for JWT storage, search cache, and Celery broker
Background Workers	Celery workers (concurrency 2), dedicated indexing queue, result backend in Redis
LLM Inference	Groq API (llama-3.3-70b-versatile), server-side only, key in environment variable
Web Search	Tavily Search API, server-side only, graceful degradation if unavailable
Embeddings	sentence-transformers/all-MiniLM-L6-v2 via HuggingFace Inference API
Monitoring	Request logging middleware, /health endpoint, Vercel Analytics, optional Sentry
Security	JWT auth (HS256, 30-min access / 7-day refresh), HTTPS, Pydantic validation, no secrets in repo

Table 5.6: CoWriteIA Production Deployment Summary

Chapter 6

Conclusions and Future Work

6.1 Conclusion

CoWriteIA was conceived to address a concrete and underserved problem: long-form writers working on complex projects have no single tool that combines project-level semantic memory, context-aware generation, character management, and research support in one coherent workflow. Existing solutions either provide isolated AI generation without narrative awareness, or strong organisational structure without any intelligence. This gap motivated the design and implementation of CoWriteIA across four development iterations.

6.1.1 What Was Built

The system delivers a full-stack, AI-assisted writing platform built on a FastAPI backend, a Next.js frontend, MongoDB Atlas as the primary document store, and ChromaDB as the vector store. The core of the platform is a Retrieval-Augmented Generation (RAG) pipeline that indexes all project documents into dense semantic embeddings using `sentence-transformers/all-MiniLM-L6-v2`, stores them in ChromaDB with a MongoDB fallback, and retrieves the most contextually relevant chunks at query time to ground every AI response in the writer's own material.

On top of this foundation, seven functional modules were implemented and integrated:

- **Project Indexing and Semantic Memory** — automatic chunking, embedding generation, and dual-store persistence (ChromaDB + MongoDB Atlas Vector Search) with background Celery workers for asynchronous indexing.
- **Context-Aware Writing Assistant** — an inline copilot that assembles RAG context, analyses writing style (POV, tense, formality), and generates suggestions that match the author's established voice.
- **Character and Scene Management** — persistent character profiles with structured attribute dictionaries, scene detection, and a relationship graph built from

co-occurrence analysis across the project corpus.

- **Dialogue Generation** — character-aware dialogue suggestions that inject personality traits, emotional tendencies, and narrative role into the LLM prompt, ensuring generated dialogue is consistent with established character voices.
- **Research Integration Agent** — real-time web search via the Tavily API, synthesised by the Groq LLM into cited, writer-focused answers with inline source attribution, persisted to conversation history.
- **Project Query Interface** — a conversational interface for asking natural-language questions about the project, answered using RAG-assembled context from the indexed corpus.
- **Gap Analysis Module** — automated identification of incomplete or underdeveloped sections, with targeted suggestions for improving narrative coherence and coverage.

6.1.2 Testing and Validation

The system was validated through a comprehensive testing strategy spanning 231 back-end tests across integration, unit, and service layers, all passing at a 100% rate. Black box testing confirmed correct API behaviour for valid and invalid inputs across all major endpoints. Unit tests isolated service logic using mocks and stubs, covering authentication, file processing, embedding generation, RAG context assembly, LLM orchestration, entity extraction, relationship discovery, edit proposals, and all three features introduced in Iteration III. Code coverage reached 32% overall, with high-coverage modules including the research endpoint (100%), web search service (92%), copilot service (87%), and edit proposal service (85%). Lower coverage in background task modules reflects their dependency on live Celery and Redis infrastructure rather than gaps in logic coverage.

6.1.3 Deployment

The final iteration brought the system to a production-ready state. The Next.js frontend is deployed on Vercel with automatic CI/CD on every push to main. The FastAPI back-end runs behind a reverse proxy with four Uvicorn workers, HTTPS enforcement, and a hardened middleware stack covering authentication, error handling, request logging, and CORS. MongoDB Atlas provides the persistent document store with continuous backups and an Atlas Vector Search index for fallback retrieval. ChromaDB runs as a co-located persistent server for primary vector search. Redis handles JWT refresh token storage, search result caching, and Celery task brokering. All secrets are managed through environment variables and never committed to the repository.

6.1.4 Reflection on Objectives

The original problem statement identified three core deficiencies in existing tools: lack of project-level context, absence of semantic memory, and disconnected workflows requiring

writers to switch between multiple applications. CoWriteIA addresses all three. The RAG pipeline provides persistent semantic memory that scales with the project. The unified workspace integrates writing, research, character management, and query support in a single interface. The context-aware generation ensures that AI output is grounded in the writer's own material rather than generic training data. The result is a platform that meaningfully reduces cognitive overhead and supports narrative consistency throughout the creative process.

6.2 Future Work

While the current system is complete and deployed, several directions could extend its capabilities and robustness.

6.2.1 Retrieval Quality and Scalability

The current RAG pipeline uses a fixed chunk size and a single embedding model. For very large projects (hundreds of thousands of words), retrieval quality could be improved through:

- **Hierarchical indexing** — maintaining both fine-grained sentence-level chunks and coarser paragraph-level summaries, allowing the retriever to operate at the appropriate granularity depending on the query type.
- **Re-ranking** — adding a cross-encoder re-ranking step after initial vector retrieval to improve precision before passing context to the LLM.
- **Hybrid retrieval** — combining dense vector search with sparse BM25 keyword search for queries that contain rare proper nouns (character names, place names) that may not be well-represented in the embedding space.

6.2.2 Evaluation Framework

The current validation relies on functional testing. A dedicated evaluation framework would strengthen confidence in generation quality:

- **RAG faithfulness metrics** — measuring whether generated responses are grounded in the retrieved context using automated metrics such as RAGAS faithfulness and answer relevance scores.
- **User studies** — structured evaluation with actual writers to assess perceived consistency, usefulness of suggestions, and reduction in context-switching compared to existing tools.
- **Style consistency scoring** — automated measurement of stylistic alignment between generated text and the author's corpus using readability and stylometric features.

6.2.3 Model and Provider Flexibility

The system currently depends on Groq (Llama 3.3 70B) for LLM inference and HuggingFace for embeddings. Future work could introduce:

- A provider abstraction layer allowing writers to select their preferred LLM (OpenAI GPT-4o, Anthropic Claude, local Ollama models) without changes to the service logic.
- Fine-tuned embedding models trained on literary corpora to improve semantic similarity for creative writing domains, where standard sentence-transformer models may underperform on stylistic or thematic queries.
- Streaming responses for the chat and copilot interfaces, reducing perceived latency for long-form generation.

6.2.4 Collaborative Writing

The current system is single-user per project. Extending it to support collaborative writing would require:

- Real-time co-editing with operational transformation or CRDT-based conflict resolution.
- Per-user contribution tracking and attribution within the shared project corpus.
- Role-based access control (author, editor, reviewer) with corresponding permission scopes.

6.2.5 Interface and Accessibility

The frontend is functional but has room for polish. Planned improvements include a richer text editor with inline diff visualisation for edit proposals, a visual character relationship graph rendered from the discovered entity network, and improved accessibility compliance across all interactive components.

6.2.6 Infrastructure Hardening

On the deployment side, future iterations could introduce container orchestration via Kubernetes for automatic scaling and self-healing, a dedicated secrets management service (AWS Secrets Manager or HashiCorp Vault) to replace environment variable injection, and a full observability stack (Prometheus metrics, Grafana dashboards, Sentry error tracking) for production monitoring beyond the current request-logging baseline.

Bibliography

- [1] M. Abdin, J. Aneja, H. Awadallah, A. Awadalla, A. Awan, N. Bach, et al. Phi-3 technical report: A highly capable language model locally on your phone. *arXiv preprint*, 2024.
- [2] Grammarly. Features. Grammarly Website, n.d.
- [3] Kelvin Guu, Kenton Lee, Zora Tung, Panupong Pasupat, and Ming-Wei Chang. Realm: Retrieval-augmented language model pre-training. *Proceedings of the 37th International Conference on Machine Learning*, pages 3929–3938, 2020.
- [4] Omar Khattab and Matei Zaharia. Colbert: Efficient and effective passage search via contextualized late interaction over bert. *Proceedings of the 43rd International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 39–48, 2020.
- [5] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive nlp tasks. *Advances in Neural Information Processing Systems*, 33:9459–9474, 2020.
- [6] Literature and Latte. Scrivener overview. Literature and Latte Website, n.d.
- [7] Notion. Everything you can do with notion ai. Notion Website, n.d.
- [8] Nils Reimers and Iryna Gurevych. Sentence-bert: Sentence embeddings using siamese bert-networks. *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*, pages 3982–3992, 2019.

Appendix A

Appendices

A.1 Appendix A: System Diagrams

A.1.1 Use Case Diagram




Figure A.1: Use Case Diagram CoWriteIA

A.1.2 Activity Diagram




Figure A.2: Activity Diagram

A.2 Appendix B: Architecture Diagram

A.2.1 CoWriteIA Detailed Architecture Diagram




Figure A.3: CoWriteIA Multi-Tier Architecture Diagram

A.3 Appendix C: User Interface Screenshots

A.3.1 Landing Page




Figure A.4: CoWriteIA Landing Page

A.3.2 Dashboard Interface




Figure A.5: CoWriteIA Dashboard Interface

A.3.3 Writing Environment Interface




Figure A.6: Writing Environment Interface

A.3.4 Writing Assistant Interface




Figure A.7: Writing Assistant Interface

A.3.5 Login Interface




Figure A.8: Login Interface

A.4 Appendix D: Detailed Use Case Specification

A.4.1 UC-01: Generate Context-Aware Writing

Use Case ID	UC-01
Use Case Name	Generate Context-Aware Writing
Primary Actor	Writer (User)
Secondary Actors	AI Processing Engine, Semantic Retrieval Module
Preconditions	• User is authenticated and logged into the system. • A project exists with at least one uploaded or created document. • Project content has been indexed and embeddings are stored.
Postconditions	• AI-generated content is displayed to the user. • Generated text may be saved to the project document.
Trigger	User selects the “Generate Content” option within the writing interface.
Main Flow	1. User opens a project document in the writing environment. 2. User requests content generation (e.g., continue writing, rewrite, expand). 3. System retrieves relevant context using semantic search. 4. System sends the prompt + context to the AI model. 5. AI model generates new content. 6. System displays the generated text to the user. 7. User may choose to accept, edit, or regenerate the content.
Alternate Flow	AF-1 User requests generation with no specific prompt. (a) System uses recent document content as context. (b) AI generates a default continuation. AF-2 User modifies the context manually. (a) User provides custom text or notes. (b) System merges custom context with retrieved chunks.
Exceptions	E1 No embeddings found: System prompts user to index project files. E2 AI service unavailable: System displays fallback message and suggests retrying. E3 Empty or invalid user request: User is asked to refine or re-enter the prompt.
Priority	High
Frequency of Use	Very frequent (core functionality)

Table A.1: Detailed Use Case: Generate Context-Aware Writing

A.5 Appendix E: Pseudo-API Snippets (Call Signatures)

- `TextExtractionService.chunk_text(file_id: str, project_id: str, text: str, start_position: int = 0, start_chunk_index: int = 0) -> List[TextChunk]`
- `EmbeddingService.generate_single_embedding(text: str) -> List[float]`
- `EmbeddingService.generate_embeddings(texts: List[str]) -> List[List[float]]`
- `EmbeddingService.store_text_chunks_with_embeddings(file_id: str, project_id: str, chunk_data: List[Dict]) -> List[TextChunk]`
- `ChromaService.add_embeddings(project_id: str, entries: List[Dict]) -> None`
- `ChromaService.search_similar(project_id: str, query_embedding: List[float], n_results: int) -> List[Dict]`
- `RAGContextService.assemble_context(query: str, project_id: str, file_id: Optional[str] = None, max_chunks: int = 5) -> Dict`
- `AIChatService.chat(request: ChatRequest, user_id: str) -> ChatResponse`